

AI Data Platform GA – November 2025

1. Introduction

AI is going to transform entire industries, redefining how organizations operate, innovate, and compete. Yet most enterprises remain unprepared for this shift. Their data is fragmented across applications, clouds, and business units; their teams rely on isolated analytics tools; and their AI initiatives struggle to reach production due to challenges in data access, governance, and model integration. To truly benefit from AI, organizations must rethink their approach and build systems that allow AI to understand their business, their processes, and their data.

So how can organizations unify their data, accelerate AI adoption, and scale intelligence across the business without adding operational overhead?

AI DATA PLATFORM

The answer is a modern AI Data Platform: a unified, end-to-end environment that seamlessly integrates data ingestion, transformation, governance, analytics, machine learning, and generative AI capabilities. It empowers organizations to work with any data—structured or unstructured, batch or real-time, operational or analytical—while providing the compute power necessary to build, deploy, and operationalize AI at scale.

The AI Data Platform fits naturally into the enterprise technology ecosystem. It consolidates data lifecycle workflows, enables governed self-service access, and makes it easy to leverage advanced analytics and AI.

Powerful yet simple, the AI Data Platform helps organizations turn their data into trusted insights and deploy AI solutions quickly and responsibly. It ensures that teams—from data engineers to business analysts to data scientists—can collaborate in a governed and efficient environment while maintaining the highest levels of performance, transparency, and control.

2. AI Data Platform Overview


The AI Data Platform is built on an open lakehouse architecture that supports structured, semi-structured, and unstructured data wherever it resides. By embracing open data formats, the platform ensures flexibility and long-term scalability for enterprise data strategies.

At its core is a unified **data and AI catalog** that inventories all enterprise assets—not only datasets, but also models, embeddings, vector indexes, and AI knowledge bases. This catalog drives end-to-end governance, security, *lineage [roadmap CY2026]*, and access control, ensuring that every analytical and AI workload operates on trusted and compliant data.

The platform provides organizations with the freedom to choose the right processing engine for each workload, whether that is Oracle Database, open-source engines such as Apache Spark or specialized processing for streaming, batch, or interactive analytics.

Surrounding the data foundation is a rich ecosystem of **AI and machine learning capabilities**, including AI foundational models, *AI Agent Studio [roadmap CY2026]*, and support for open-source model libraries. Data scientists can bring their preferred models and frameworks and apply them directly to enterprise data.

All of this is delivered through an **AI-powered developer experience** that supports the full spectrum of personas: data engineers building pipelines, data stewards managing governance, data scientists training models, analysts exploring insights, and AI developers creating intelligent agents and applications. The AI Data Platform unifies these capabilities into a single environment to accelerate the development, deployment, and operationalization of Data and AI across the enterprise.



3. Key Business Benefits

- **Unified Data Foundation**
Consolidates structured, semi-structured, unstructured, and streaming data into a single, governed platform to ensure consistency, reliability, and trust across all analytics and AI initiatives.
- **Accelerated AI Development**
Provides ready-to-use AI foundational models, vector indexing, and retrieval pipelines, enabling faster building and deployment of AI agents and applications.
- **Reduced Time-to-Insight**
Integrates data preparation, analytics, ML, and AI into one environment, eliminating silos and dramatically shortening the path from raw data to meaningful insights.
- **Lower Total Cost of Ownership (TCO)**
Simplifies operations by replacing fragmented tools with a unified platform for data engineering, governance, analytics, and AI—reducing integration, maintenance, and infrastructure overhead.
- **Open and Flexible Architecture**
Supports open data formats, open-source engines, custom model libraries, and multi-cloud interoperability, giving organizations full choice over technologies and future direction.
- **Scalable and Elastic Processing**
Adapts seamlessly to enterprise workloads—batch, streaming, interactive analytics, and AI inference—scaling up or down based on demand to optimize performance and cost.
- **Enterprise-Grade Governance & Security**
Centralized catalog-driven governance ensures consistent access controls, lineage *[roadmap CY2026]*, model transparency, data policies, and compliance across data and AI workloads.
- **Enhanced Productivity for All Personas**
Offers AI-powered experiences and tools for data engineers, [data stewards](#), data scientists, analysts, and AI developers—improving collaboration and accelerating project delivery.
- **Faster Innovation & Business Agility**
Empowers teams to rapidly prototype, iterate, and productionize both data pipelines and AI solutions, enabling organizations to innovate quickly and adapt to market demands.

By combining open formats, elastic storage, flexible processing engines, and autonomous high-performance serving layers, the AI Data Platform delivers an enterprise-grade lakehouse foundation that supports both modern analytics and next-generation AI applications.


4.2 Unified Catalog and Governance

The AI Data Platform includes a **centralized master catalog** that unifies all enterprise data and AI assets—structured, semi-structured, and unstructured—into a single, discoverable system of record. This catalog allows teams to easily find, organize, classify, and govern datasets, models, embeddings, vector indexes, and other AI artifacts.

The catalog natively supports the **medallion architecture**, providing full visibility and lineage *[roadmap CY2026]* across bronze, silver, and gold layers. This ensures that data engineers, analysts, and AI developers can trace how data was ingested, transformed, and prepared for analytics and AI workloads.

With **multi-format data support**, the platform seamlessly catalogs and manages assets from diverse sources, including databases, files, documents, images, and application logs, offering a holistic view of all enterprise data.

Built-in governance controls ensure secure and compliant data operations. Access policies, permissions, and security rules are centrally defined and enforced across every dataset and AI asset—enabling consistent, governed access for all personas and workloads.



Together, these capabilities deliver a trusted, governed foundation for analytics, machine learning, and AI applications across the enterprise.

4.3 Developer Experience

The AI Data Platform provides a unified **developer workbench** that brings all data and AI workloads together in a single environment. Developers, data engineers, data scientists, and AI Agents builders can design, integrate, process, and operationalize their solutions without switching tools or environments.

Key capabilities include:

- **Unified IDE for End-to-End Development**

A single environment that supports data integration, data engineering, data science, analytics, and AI agent development. Offers both **low-code** visual design tools and **pro-code** experiences for advanced developers.

- **AI-Assisted Notebook Interface**

Multi-language notebooks supporting **SQL, Python, Scala and Java** [roadmap CY2026], powered by AI assistance to help generate code, optimize queries, and accelerate development.

- **Flexible Compute Options**

Developers can run workloads on their choice of processing engines, including **Apache Spark, Autonomous AI Lakehouse**, or a combination of both—optimized for batch, streaming, or interactive execution.


- **Multi-Language & Team Collaboration**

Integrated **Git support** [roadmap CY2026] enables version control, collaboration, and CI/CD workflows across teams.

- **Unified Workflow Orchestration**

End-to-end pipeline orchestration to schedule, automate, and monitor jobs across data ingestion, transformation, ML training, and AI agent workflows.

Together, these capabilities provide a powerful and intuitive developer experience that streamlines the design, deployment, and operationalization of data and AI solutions across the enterprise.



4.4 AI Models and Framework

The AI Data Platform provides a comprehensive set of **AI models, frameworks, and agent tools** that enable organizations to build intelligent, domain-specific applications on top of their enterprise data.

- **Foundation Models**

Supports industry-leading AI foundational models—including **OpenAI, Grok, Llama, Cohere**, and soon **Gemini**—for tasks such as summarization, conversational AI, semantic search, and knowledge extraction.

- **Agent Studio Integration** *[roadmap CY2026]*

A built-in Agent Studio within the developer workbench enables **low-code and pro-code** creation of AI agents for workflow automation, document understanding, business monitoring, and domain-specific reasoning.

- **Vectorized Knowledge Base** *[roadmap CY2026]*


Integrated with the **Oracle 26ai Vector Store** to build high-performance RAG pipelines and semantic search over unstructured and structured enterprise data.

- **Popular AI Frameworks**

Provides support for widely used AI libraries, including **PyTorch, TensorFlow, LangChain, LangGraph**, and more—giving developers the freedom to use their preferred tools.

- **Optimized Compute & Lifecycle Management**

Models and agents run on optimized **OCI CPU and GPU shapes**, with full lifecycle management for AI assets across development, deployment, and governance. This unified model and framework ecosystem enables organizations to rapidly develop AI workflows, analytics, and domain-specific agents that automate tasks, unlock insights, and transform enterprise productivity.



5. Enterprise-Grade Security

The AI Data Platform provides secure, governed access to data and AI assets by combining OCI Identity and Access Management (IAM), workspace-level isolation, centralized catalog governance, and comprehensive audit capabilities. Security principles are consistent across data, notebooks, workflows, and collaboration assets, ensuring controlled and compliant usage.

Key Capabilities

Identity & Access Management

Access is managed through OCI IAM, including OCI policies, sign-on policies, and support for federated identities. IAM governs who can access or administer AI Data Platform resources in the OCI Console.

Role-Based Access Control (RBAC)

In addition to OCI IAM, the platform provides fine-grained RBAC for workspaces, catalogs, datasets, and other assets, enabling organizations to define precise permissions for viewing, editing, or managing resources.

Workspace Isolation

Workspaces serve as isolated logical environments that securely group notebooks, workflows, data connections, and libraries, enabling controlled collaboration across teams and projects.

Centralized Catalog Governance

The Master Catalog centralizes metadata management, tagging, and asset organization across structured and unstructured sources. It enforces catalog-level access control and supports secure sharing using Delta Sharing.

Private Network Deployment


Workspaces can be deployed inside an OCI Virtual Cloud Network (VCN) using private subnets and Network Security Groups (NSGs), enabling secure access to private data sources not exposed publicly.

Audit Logging

AI Data Platform maintains detailed audit logs of user activities, resource modifications, and workspace operations, supporting compliance monitoring and operational visibility.

OCI Object Storage Foundation

AI Data Platform is built on OCI Object Storage as its primary data layer, leveraging Object Storage's native access controls and secure, managed infrastructure for storing and accessing data across the platform.



6. Pricing

Oracle AI Data Platform uses a transparent, consumption-based pricing model that follows standard OCI principles. Customers pay only for the resources they use, including compute, storage, and platform capabilities. This provides predictable cost management and eliminates the overhead of managing multiple vendor contracts or fragmented billing sources.

A key advantage of AI Data Platform is that Oracle provides both the AI platform and the compute infrastructure under a single operational and billing model. All components—workspaces, compute clusters, notebooks, Spark jobs, **Autonomous AI Lakehouse** usage, and Object Storage—are managed and controlled from a single console. This unified approach simplifies cost visibility, governance, and capacity planning for enterprises.

Key Pricing Characteristics

- Unified Platform + Compute Model**
 AI workloads run on OCI compute and storage directly managed by the platform, providing a single point of control for provisioning, scaling, and monitoring resource usage.
- Pay-as-You-Go Consumption**
 Customers pay only for the compute, storage, and services they consume. There are no upfront commitments required.
- Consistent with OCI Pricing**
 AI Data Platform leverages OCI's globally recognized cost-efficient

infrastructure, benefiting from transparent and competitive pricing across compute, storage, networking, and AI services.

- **Full Cost Visibility**

All resource usage—Spark clusters, **Autonomous AI Lakehouse** sessions, Object Storage, workflows, and notebooks—is tracked within OCI, enabling clear, unified cost reporting and budgeting.

- **Operational Efficiency**

Because the platform and compute are integrated, organizations avoid the complexity and cost typically associated with stitching together multiple third-party tools or managing separate AI and data systems.

Oracle’s unified architecture enables customers to run data engineering, analytics, and AI workloads at scale with **lower operational overhead and highly competitive pricing**, while maintaining centralized governance and full transparency into platform usage.

Edit f1_demo_cluster

Cluster configuration
☐ Quickstart ☒ Custom
The presence of an installed library prevents a quick start, the compute will perform a standard start

AIDP Unit
546 AIDP units per hour

Driver
Driver shape
AMD

OCPUs
2

Memory (GB)
16

Block volume (GB)
200

Worker
Worker shape
AMD

OCPUs
2

Memory (GB)
16

Block volume (GB)
200

Set number of workers
☒ Static amount ☐ Autoscale

Number of Workers
2

Run duration
☐ Forever ☒ Idle timeout

Idle timeout in minutes
180


Advanced options - Adding Spark configurations and Init script will result in slower startup

Cancel

Save and Restart

11. Summary

Oracle's AI Data Platform provides a unified, open, and governed foundation for accelerating enterprise data and AI initiatives. Built on an open lakehouse architecture with a centralized catalog, integrated developer experience, and industry-leading AI models, the platform streamlines the entire lifecycle—from data ingestion and transformation to analytics, workflows, and AI agent development. With enterprise-grade security, transparent consumption-based pricing, and the simplicity of running data and AI on a single, integrated OCI platform, organizations can innovate faster, reduce complexity, and maximize the value of their data and AI investments.



Additional RFP Topics

- Networking
- High Availability (HA)
- Disaster Recovery (DR)
- Support SLAs
- Deployment Options
- CI/CD
- MLOps
- DevOps
- Onboarding Methodology
- Training & Enablement
- Product Roadmap