



Oracle Converged Database

Developing document-oriented applications
with JSON, REST and SODA

Witold Swierzy

Oracle EMEA Data Domain Expert

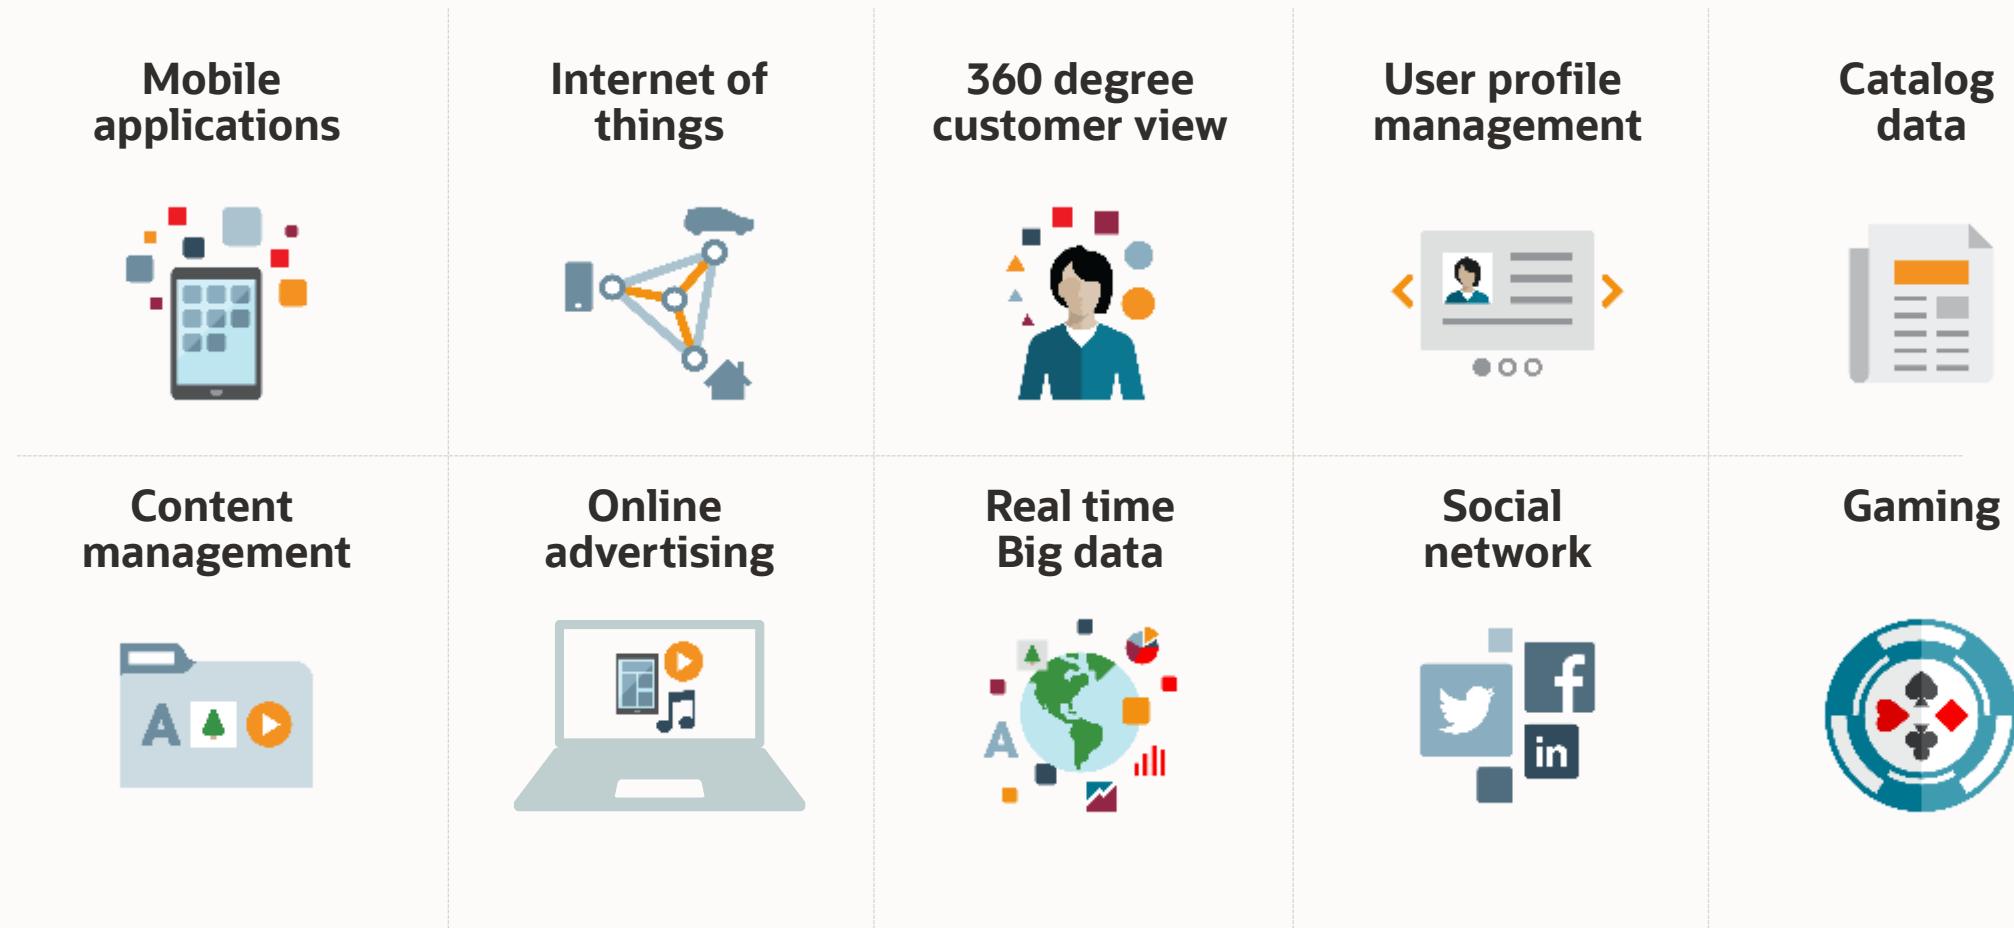
Agenda

- **Introduction**
- **Oracle Document-Oriented Database**
- **Summary**
- **SODA Live Demo**
- **Q&A**



Introduction

Modern Applications need flexible schema



Why JSON can be an answer ?

1: JSON easily maps to Business Objects

- no need to map object to multiple tables
- single insert/query/update
- less complex SQL



Customer: Major US Insurance:

- every insurance policy is a JSON doc
- different policy shapes
- all policy information in one doc
- no decomposition/joins
- low latency puts/gets
- great for web-application
- Relational views
 - reporting
 - daily risk analysis



2: JSON is schema flexible

- adding additional fields
 - *add a middle name or birth name*
- cardinality change
 - *1 address -> multiple addresses*
- Productivity, Agility

```
{  
  "firstName" : "John",  
  "middleName": "Jeremy",  
  "lastName"  : "Smith",  
  "age":25,  
  "address": [ { ... }, { ... } ]}
```

Customer: Major Global Stock Exchange

- metadata for financial products
- 1600 possible attributes
- 100 used on average
- application adds new attributes if needed without database op
- attributes in JSON column
- Custom reporting for banks
- SQL/JSON operations



3: JSON Document Store APIs

- Avoid SQL for simple (CRUD) operations
 - insert/find/update/delete a person
 - noSQL-style document APIs

```
SODA get -f {"lastName": "Smith"}
```

versus

```
SELECT id, data  
FROM persons  
WHERE data.lastName = 'Smith'
```

- Preserve SQL for reporting, analytics, ML

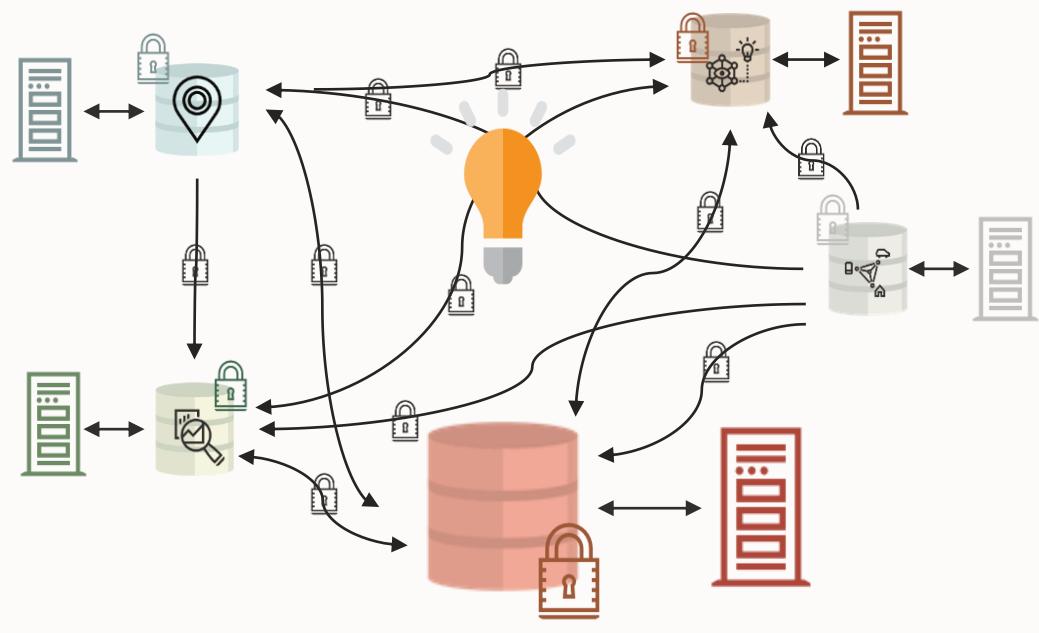
Customer: Global Retailer, stores and online shop

- Point of sale application
 - order status tracking
 - payment tracking
 - gift card processing
- Developers preferred document store model
 - agile, productive, easy to use
- DBA and Ops preferred Oracle
 - transactions
 - SQL for reports
 - security features,...

NoSQL

SQL

Multiple single-purpose databases may drive to complex and unmanageable architecture



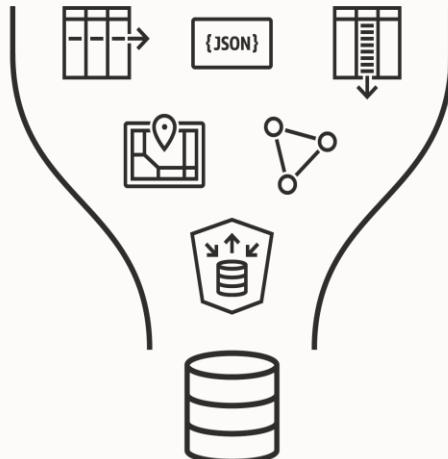
→ ETL, replication, events etc.

Traditional paradigm assumes using many single-purpose databases drives to

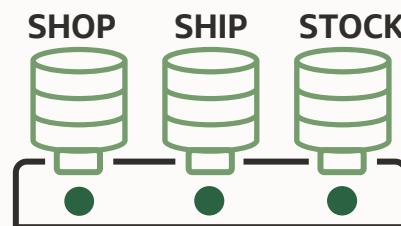
- High level of complexity of maintenance and development process
- High TCO Costs

Oracle Database

A Converged, Open SQL Database



Converged Database



Support for microservices

Multi-model

Best-of-Breed Relational, JSON, Spatial, Graph, Cube, Text, Blockchain
Cross-model operations enables you to easily create value across all your data

Multi-workload

High Performance Transactions, DW, Analytics, ML, IoT, Streaming, Multitenant
Deep optimizations deliver exceptional price-performance across all workloads

Most productive for developers and analysts

Same SQL and transactions operate on any data and workload
Integrated microservices, events, REST, CI/CD, Low-code

Support for microservices

Low-level data types and workloads should not dictate your architecture

Oracle Document-Oriented Database



Oracle Converged Database as a **{JSON}** Document Database

support for all modern Languages/Drivers/Tools/APIs



Oracle Converged Database as a **{JSON}** Document Database

cross-model DBMS allowing for holistic view of all the data

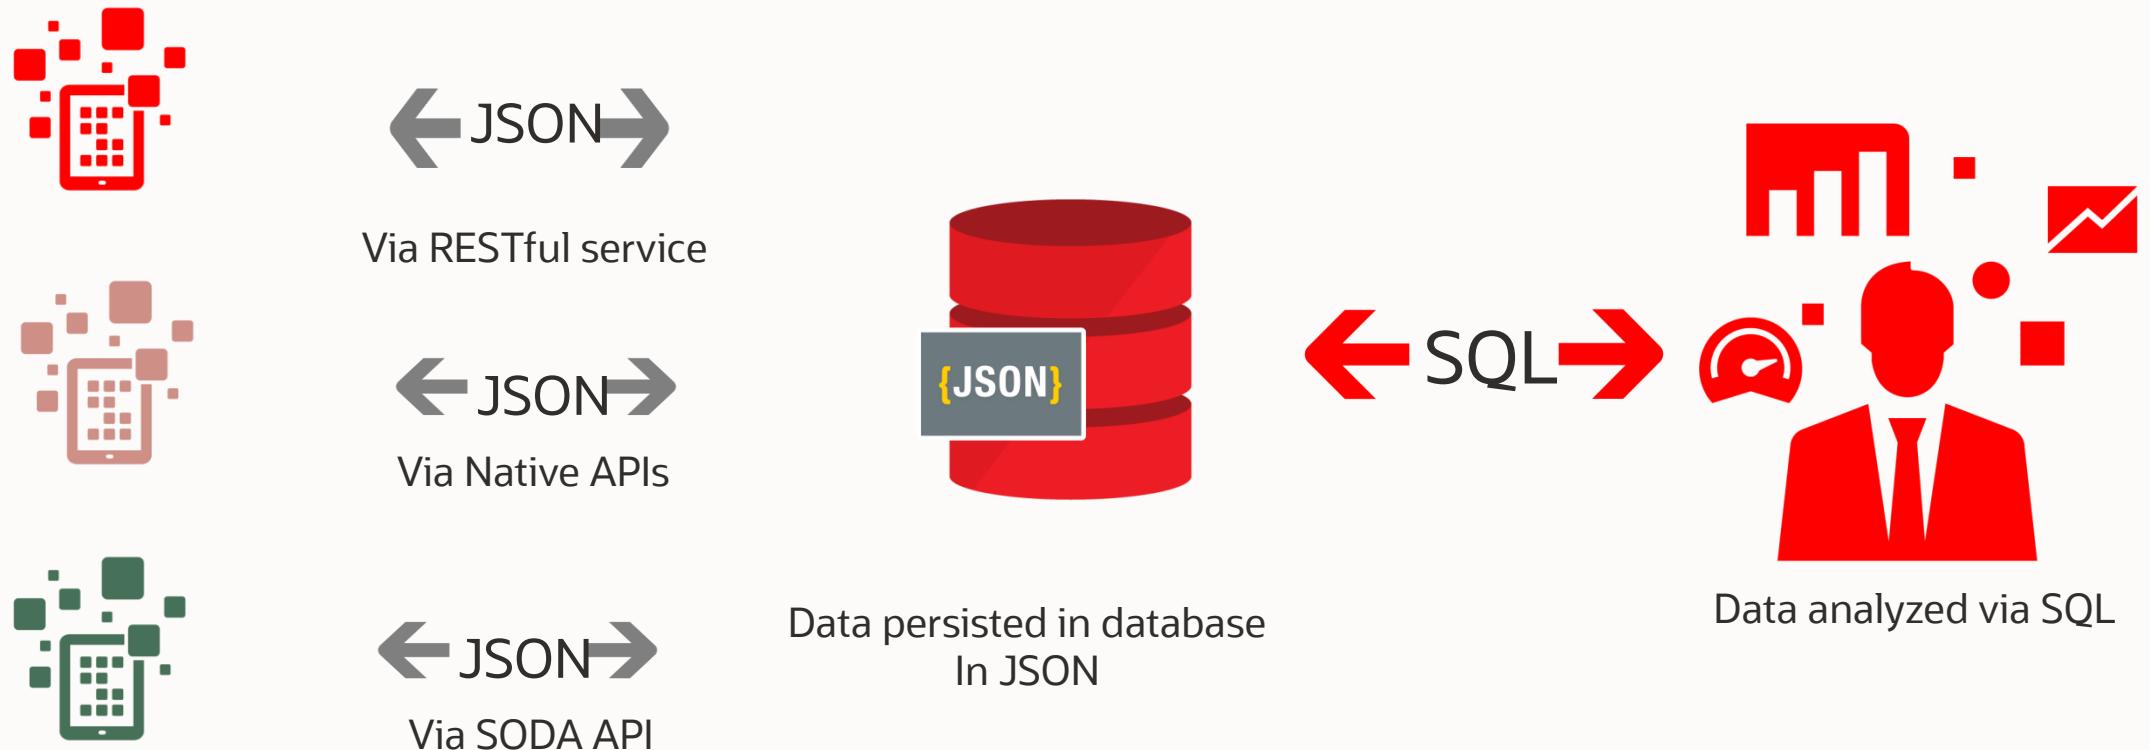


```
select p.id, p.name,  
       sl.json_doc.quantity Anzahl  
  from  
product p, customer c, store s, sales sl  
 where  
sl.json_value(json_doc,'$.product_id')=p.id  
and  
contains(p.notes,fuzzy('Lifferkosten'), 1) > 0  
and  
sl.json_value(json_doc,'$.CUSTNumber'  
returning number) = c.cust_ids  
and  
s.store_id = 1234  
and  
sdo_within_distance(c.location,  
s.location,'distance=20')='TRUE'
```

Oracle JSON

Oracle Converged Database as a **{JSON}** Document Database

Access JSON data via multiple API(s)



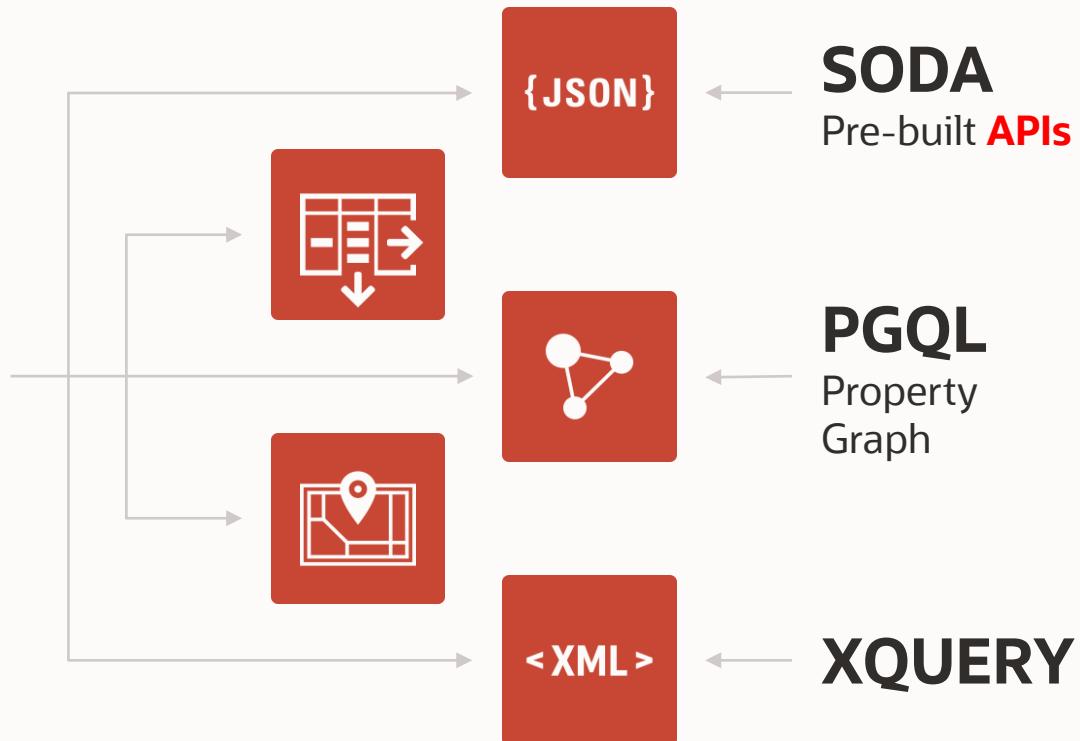
Oracle Converged Database as a {JSON} Document Database

Access data via SQL & REST or modern NoSQL APIs: your choice

Cross-Model Data Access

SQL & REST

Relational, Graph,
Document, Spatial,
Temporal, multidimensional



NoSQL Data Access

SODA

Pre-built APIs for JSON

PGQL

Property Graph

SPARQL

RDF Graph

XQUERY

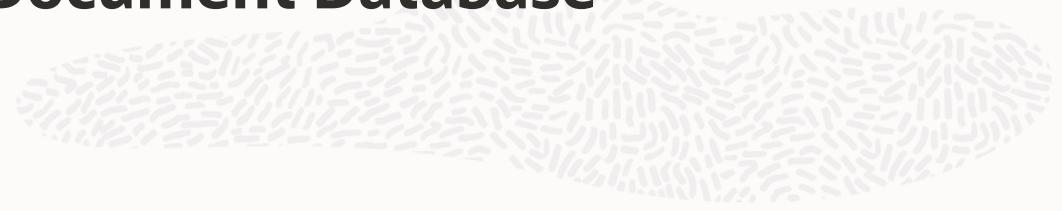
Developers can easily access multiple data models simultaneously via SQL, REST and APIs

Oracle Converged Database as a **{JSON}** Document Database

- Full SQL support
- ACID transactions
- JSON search indexes
- Advanced security
- APEX low-code development



Oracle Converged Database as a **{JSON}** Document Database



- You have the choice to store JSON data as
 - Collections
 - Column(s) in relational tables
- JSON documents are stored as
 - Native Binary data type
 - **JSON** (up to 32 MB)
 - Textual data type
 - VARCHAR2 (up to 32* KB)
 - BLOB (up to 2 GB)
 - CLOB (up to 1 GB)



Oracle Converged Database as a **{JSON}** Document Database

SQL support for JSON

- **SQL** for advanced reporting and analytical operations on JSON documents
 - Dot notation
 - SELECT coll.doc.address.postalCode...
 - SQL/JSON Path Expressions
 - Rich set of functions
 - Views + On-Demand Refreshable Materialized Views
 - Analytical functions (over (), window (...))
 - Documents Pattern Matching (match_recognize)
 - JSON documents generation from relational data
 - ...
- **PL/SQL** for **server-side** processing
 - Parse and generate Documents
 - apex_json
 - Secure access to JSON Open Data
 - apex_web_service

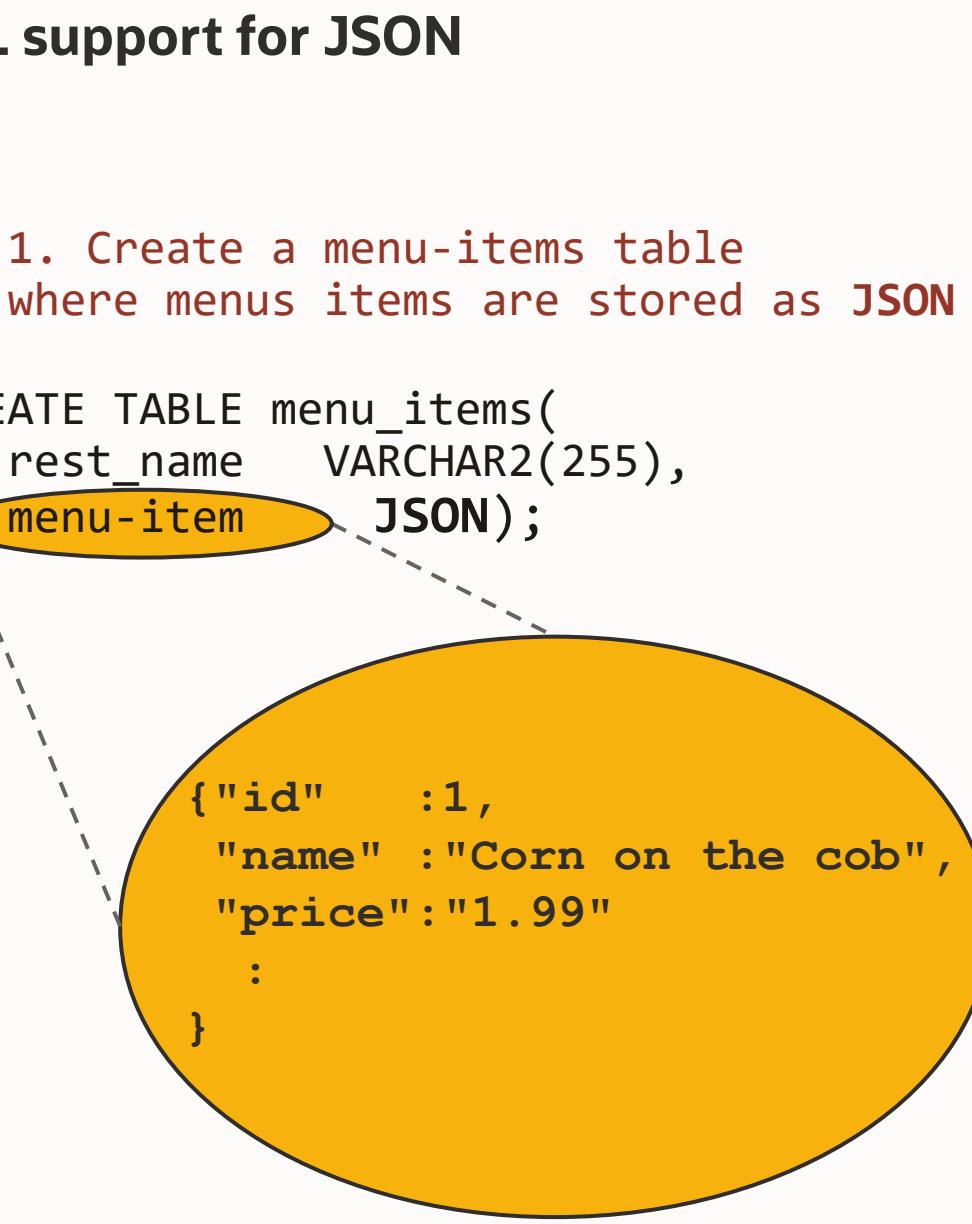


Oracle Converged Database as a {JSON} Document Database

SQL support for JSON

```
-- 1. Create a menu-items table  
-- where menus items are stored as JSON
```

```
CREATE TABLE menu_items(  
    rest_name  VARCHAR2(255),  
    menu-item  JSON);
```



```
{"id": 1,  
 "name": "Corn on the cob",  
 "price": "1.99"  
 :  
 }
```

-- 2. Use simple dot notation to access
elements within the JSON docs

```
SELECT m.menu_item.name item,  
      m.menu_item.price price  
FROM   menu_items m;
```

ITEM	PRICE
-----	-----
Corn on the cob	1.99

-- 3. Create index on a JSON column
element to speed up rows filtering

```
CREATE INDEX menu_items_idx ON  
purchase_orders  
( JSON_VALUE(JSON_DOCUMENT, $.id.number())  
));
```



Oracle Converged Database as a **{JSON}** Document Database

SQL support for JSON



Native Support to Index and Search JSON Documents

Oracle makes it simple for Apps to index, search and analyze text

Search text using keyword search, context queries, pattern matching, etc.

- Search websites, catalogs, documents, LOBs

Perform linguistic analysis on documents to easily classify them

- Classify customer feedback as positive, negative, or neutral using sentiment analysis

Search JSON or XML documents using the structure of the document to restrict the search

- Find all purchase orders where the comment field contains “damage” and “delivered”

Oracle Converged Database as a {JSON} Document Database

SQL support for JSON



Native Support to Index and Search JSON Documents

-- 1. Create a text index on the JSON column of the purchase_orders table

```
CREATE SEARCH INDEX PO_search_ind ON purchase_orders(po_doc) FOR JSON;
```

--2. Search for PO documents with “Delivered” and “Damage” in the comments field

```
SELECT po_number, po_doc.comment comment
FROM   purchase_orders
WHERE  JSON_TEXTCONTAINS(po_doc, '$.comment', 'Delivered and Damage');
```

Oracle Converged Database as a {JSON} Document Database

SQL support for JSON

Ensuring consistency

- Constraints

```
CREATE UNIQUE INDEX unique_idx ON purchase_orders
( JSON_VALUE(JSON_DOCUMENT, '$.PONumber.number()' ));  
  
SQL> soda insert purchase_orders {"PONumber": 1601 };  
SQL> commit;  
  
SQL> soda insert purchase_orders {"PONumber": 1601 };
```

ORA-00001: unique constraint (HR.UNIQUE_IDX) violated

```
ALTER TABLE purchase_orders ADD CONSTRAINT
po_address_zipcode_positive_number CHECK (
    JSON_VALUE( json_document,
        '$.ShippingInstructions.Address.zipCode'
        returning number)
    > 0 )
ENABLE      -- enable constraint
NOVALIDATE; -- but don't check existing documents
```

- Triggers

```
CREATE TRIGGER validate_items
AFTER UPDATE ON purchase_orders FOR EACH ROW
DECLARE l_number_of_products number;
BEGIN
    SELECT sum(Quantity) into l_number_of_products
    FROM JSON_TABLE(:NEW.json_document, '$'
        Columns( Nested Items[*] Columns(Quantity NUMBER)));
    IF l_number_of_products > 10 THEN
        RAISE_APPLICATION_ERROR(-20001, 'No more than 10
            products ordered at the same time.');
    END IF;
END;
/  
  
SQL> UPDATE purchase_orders SET json_document =
    JSON_TRANSFORM( json_document, APPEND '$.Items' =
    '{"Description":"Shiny toy", "UnitPrice":1.0,
        "UPCCode": 123456789, "Quantity": 5}' format json );  
  
ORA-20001: No more than 10 products ordered at the same time.
```

Oracle Converged Database

SODA (Simple Oracle Document Access) and SQLcl



SODA: Goals

- Enable schemaless development on top of an Oracle Database
 - Provide a simple NoSQL-style API for working with documents
- Make it easy to use Oracle as a NoSQL-style document store
 - Allow developers to work with Oracle without learning SQL
 - Allow developers to work with Oracle without DBA support
- Support all common application development environments
 - Traditional programming languages
 - Scripting languages and frameworks



Oracle Converged Database

SODA (Simple Oracle Document Access) and SQLcl



SODA APIs

- NoSQL-style APIs for
 - Java, JavaScript/Node.js, Python, REST, PL/SQL, C...
- Used to manage JSON data
 - create collections
 - store documents in collections
 - retrieve documents
 - query documents
- **No need to know SQL!**

SQLcl

- Modern SQL Developer Command Line interface for Oracle database
- Provides
 - inline editing, statement completion, command recall...
- **SODA commands**

Oracle Converged Database

SODA (Simple Oracle Document Access) and SQLcl



SQLcl

- Modern SQL Developer Command Line interface for Oracle Database
- Provides
 - Inline editing, statement completion, command recall...
 - **SODA commands**
 - **Full CRUD operations support**

```
SODA allows schemaless application development using the JSON data model.
SODA create <collection_name>
Create a new collection

SODA list
List all the collections

SODA get <collection_name> [-all | -f | -k | -klist] [{<key> | <k1> <k2> ... > | <qbe>}]
List documents the collection
Optional arguments:
  -all    list the keys of all docs in the collection
  -k     list docs matching the specific <key>
  -klist  list docs matching the list of keys
  -f      list docs matching the <qbe>

SODA insert <collection_name> <json_str | filename>
Insert a new document within a collection

SODA drop <collection_name>
Delete existing collection

SODA count <collection_name> [<qbe>]
Count # of docs inside collection.
Optional <qbe> returns # of matching docs

SODA replace <collection_name> <oldkey> <new_{str|doc}>
Replace one doc for another

SODA remove <collection_name> [-k | -klist | -f] {<key> | <k1> <k2> ... > | <qbe>}
Remove doc(s) from collection
Optional arguments:
  -k     remove doc in collection matching the specific <key>
  -klist remove doc in collection matching the list <key1> <key2> ... >
  -f      remove doc in collection matching <qbe>
```

Oracle SODA

examples

Node.js

```
conn = await oracledb.getConnection(...);
db = conn.get SodaDatabase();
col = await db.createCollection("purchase_orders");
await col.drop();
```

Python

```
conn = cx_Oracle.connect(...);
db = conn.get SodaDatabase();
col = db.createCollection("purchase_orders");
col.drop();
```

Java

```
OracleClient client = new OracleRDBMSClient();
db = client.getDatabase(jdbcConn);
OracleCollection col =
db.admin.createCollection("purchase_orders");
col.admin().drop();
```

PL/SQL (and Oracle Application Express)

```
col := dbms_soda.create_collection('purchase_orders');
select dbms_soda.drop_collection('purchase_orders')
from dual;
```

Oracle SODA

Sample services provide by SODA for REST



GET /DBSODA/schema	List all collections in a schema
GET /DBSODA/schema/collection	Get all objects in collection
GET /DBSODA/schema/collection/id	Get specific object in collection
PUT /DBSODA/schema/collection	Create a collection if necessary
PUT /DBSODA/schema/collection/id	Update object with id
POST /DBSODA/schema/collection	Insert object into collection
POST /DBSODA/schema/coll?action=query	Find objects matching filter in body

Oracle REST Data Services

ORDS powers the Schema Service, A Quick Detour

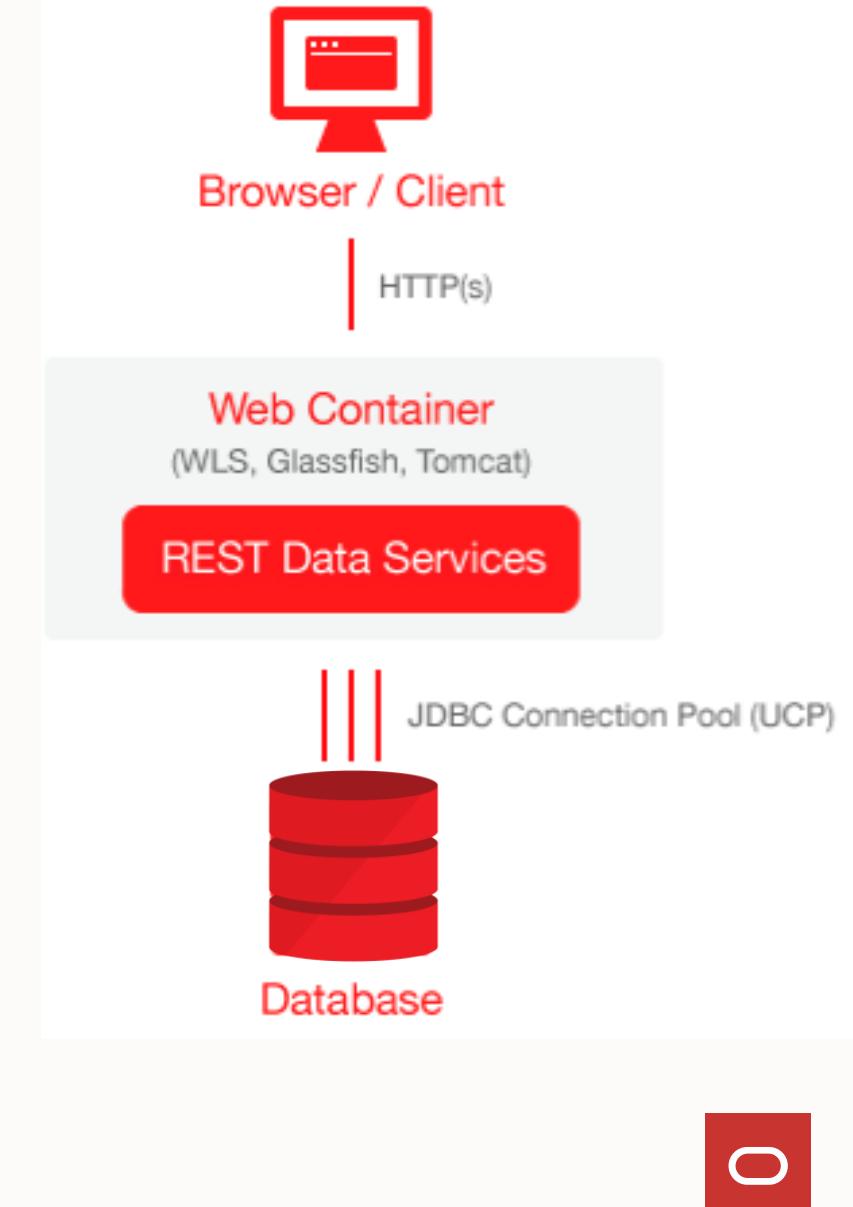
- Provides data access consistent with modern App Dev frameworks
 - Mid tier application
 - Can map standard http(s) RESTful gets and posts to SQL
 - Can declaratively returns results in JSON format
 - JavaScript friendly
 - Can support high numbers of end users
- Services
 - HTTP(s) relational data access
 - Oracle JSON collection based schema-less access
 - Oracle NoSQL access over HTTP
 - Oracle APEX mid-tier, web toolkit applications, mod_plsql replacement



Oracle Converged Database

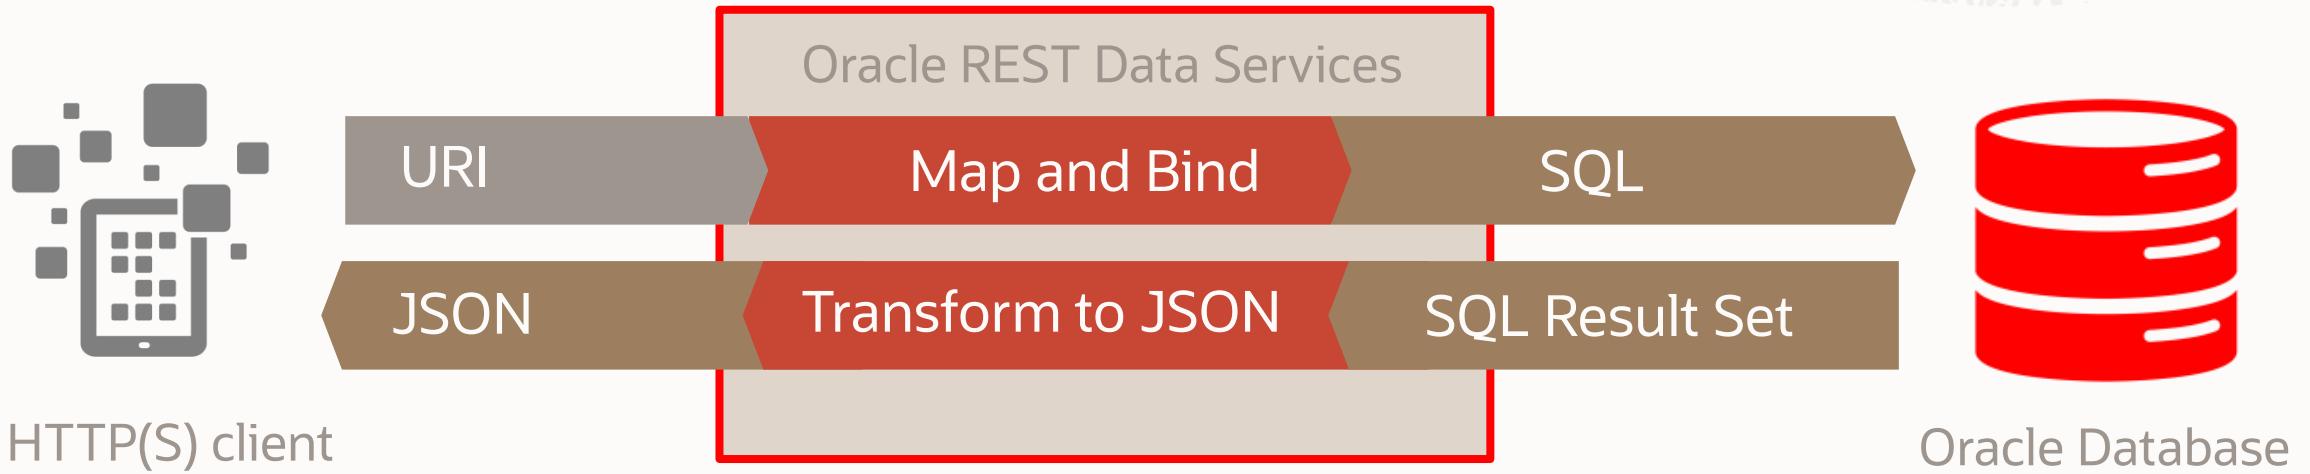
Oracle REST Data Services

- Available **Standalone (Jetty)**, Weblogic, Tomcat & Glassfish
- Turns Database Service into an RESTful API service
- Fully provisioned and functional in all cloud editions
- Available since 11g, **no extra cost**
- Allows publishing of URI based access to Oracle database over REST
- Results in JSON or CSV
- Mapping of URI to SQL or PL/SQL
- All HTML methods GET, PUT, POST, DELETE, PATCH
- OAuth2 integration
- Highly scalable, can use all features of database



Oracle REST Data Services

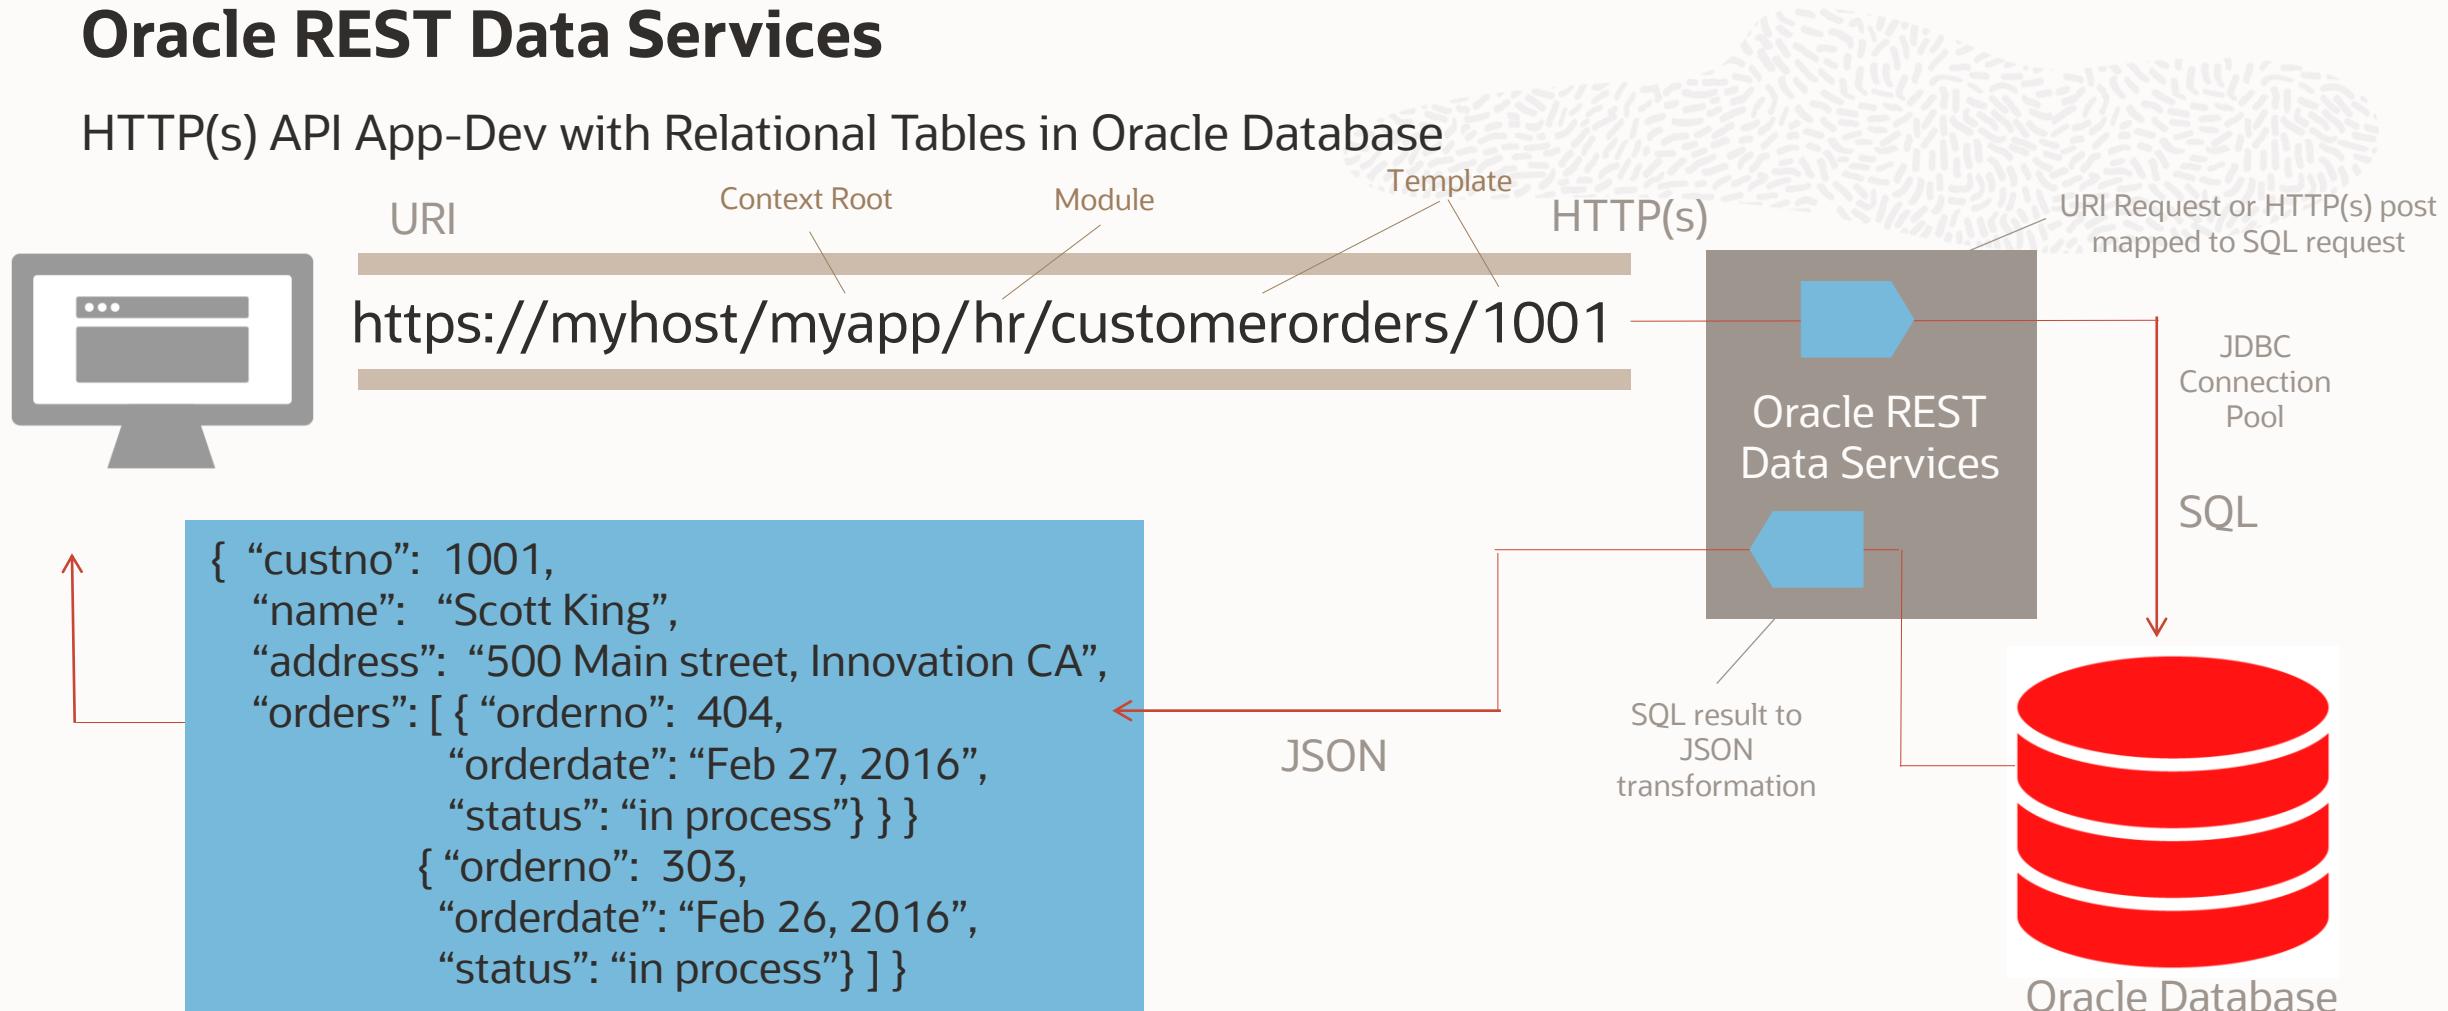
Serving JSON results from relational data



- Data stored in standard relational tables and columns
- Oracle REST Data Services (ORDS) Developer defines URI<>SQL mapping
- App Developer calls named URI over HTTP(S) gets and posts

Oracle REST Data Services

HTTP(s) API App-Dev with Relational Tables in Oracle Database



ORDS maps standard URI requests to corresponding relational SQL (not schemaless): e.g. SQL SELECT from customers and orders table.

ORDS also transforms the SQL results into JavaScript Object Notation (JSON), other formats include HTML, binary and CSV.

Fully committed to supporting any and all standards required by Fusion / SaaS / FMW; we are actively engaged in the ongoing dialog.

Oracle REST Data Services

RESTful development with JSON

RESTful services are a simple well, understood model

CRUD operations map to HTTP Verbs

Create/Update : PUT / POST
Retrieve:GET – Delete :
DELETE

Other operations, such as Query by Example, Bulk Insert and Indexing are mapped to variants of POST

JSON document forms the payload of the HTTP Request or Response

Stateless model, no transaction support

Oracle Autonomous Database - support for MongoDB

Develop and run MongoDB workloads in the Oracle Autonomous Database



Modern document-centric development

- JSON Collections-based data model
- Rich clients – MongoDB API, REST and SODA based development API
- Native JSON storage with advanced indexes and optimized performance

... and proven enterprise functionality

- ACID Transactions
- SQL-based Reporting and Analytics (including scalable parallel execution)

... running on the Autonomous Database platform

- Availability
- Security
- Elasticity

Oracle Autonomous Database - support for MongoDB

Develop and run MongoDB workloads in the Oracle Autonomous Database



API compatibility details

MongoDB API version 4.2 compatible

- Transactions supported
- Load balanced connections supported like on MongoDB Atlas

MongoDB API backed by Oracle user and roles management instead of MongoDB's

- Allows Oracle's enterprise-class security features to be used with MongoDB collections
- Makes unified user management easy

Oracle Autonomous Database - support for MongoDB

Develop and run MongoDB workloads in the Oracle Autonomous Database

JSON Development IDE

The screenshot shows the Oracle Database Actions JSON interface. A search bar at the top right contains the text "MERNDemo". Below it, a table displays a single document with the query filter `{"brand": "Apple"}`. The document details are as follows:

```
{  
  "price": 89.99,  
  "rating": 0,  
  "countInStock": 3,  
  "description": "Bluetooth technology lets you connect it with compatible devices wirelessly. High-quality AAC audio offers immersive listening experience. Built-in microphone allows you to take calls while working.",  
  "category": "Electronics",  
  "createdAt": "2022-01-14T22:30:27.528000Z",  
  "_id": "61e1f983746641a5d4af602a",  
  "updatedAt": "2022-01-14T22:30:27.528000Z",  
  "_v": 0,  
  "reviews": [],  
  "name": "Airpods Wireless Bluetooth Headphones",  
  "numReviews": 0,  
  "user": "61e1f983746641a5d4af6027",  
  "brand": "Apple",  
  "image": "/images/airpods.jpg"  
}  
  
{  
  "price": 599.99,  
  "rating": 0,  
  "countInStock": 10,  
  "description": "Introducing the iPhone 11 Pro. A transformative triple-camera system that adds tons of capability without complexity. An unprecedented leap in battery life",  
  "category": "Electronics",  
  "createdAt": "2022-01-14T22:30:27.528000Z",  
  "_id": "61e1f983746641a5d4af602b",  
  "updatedAt": "2022-01-14T22:30:27.528000Z",  
  "_v": 0,  
  "reviews": [],  
  "name": "iPhone 11 Pro 256GB Memory",  
  "numReviews": 0  
}
```

At the bottom left, there are two small icons: a triangle and a circle. At the bottom right, the status bar shows "10:33:19 PM - REST call resolved successfully."

Oracle Cloud tools

The screenshot shows the MongoDB Compass interface for the `merndemo.products` collection. The interface includes tabs for **Documents**, **Aggregations**, **Schema**, **Explain Plan**, **Indexes**, and **Validation**. The **Documents** tab is selected, displaying 7 documents. A filter is applied: `{"brand": "Apple"}`. The results show two documents matching the filter, both of which are Apple products. The first document is for Airpods Wireless Bluetooth Headphones and the second for an iPhone 11 Pro 256GB Memory.

Document	Product Name	Brand	Price
1	Airpods Wireless Bluetooth Headphones	Apple	89.99
2	iPhone 11 Pro 256GB Memory	Apple	599.99

MongoDB Compass and other tools

Oracle Database

The most advanced database platform in the market

Document-oriented data model is fully supported by all the database options

Consolidated and Mixed-Workloads



Multitenant Database, Exadata, OLTP and analytics in one

Fault-Tolerant Availability



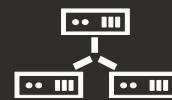
RAC, Data Guard, Flashback, Online Data Reorg, Online Patching, MAA

Highest Security



Data Safe, Advanced Security, Encryption, Roles, Privileges ...

Transparent Scaling



RAC, Exadata, Parallel SQL, Native Sharding

Comprehensive Data Model Support



JSON, text, graph, spatial, blockchain

Powerful Analytics

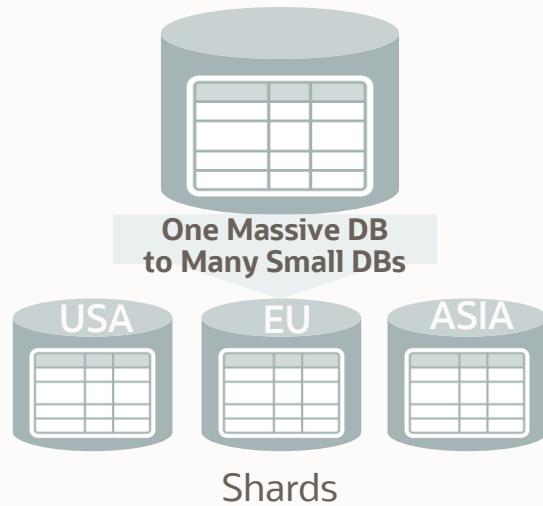


In-memory analytics, machine learning, data partitioning

These technologies fully support document-oriented Oracle Converged Database

Oracle Converged Database

Native Sharding Simplifies Distributed Data Architecture



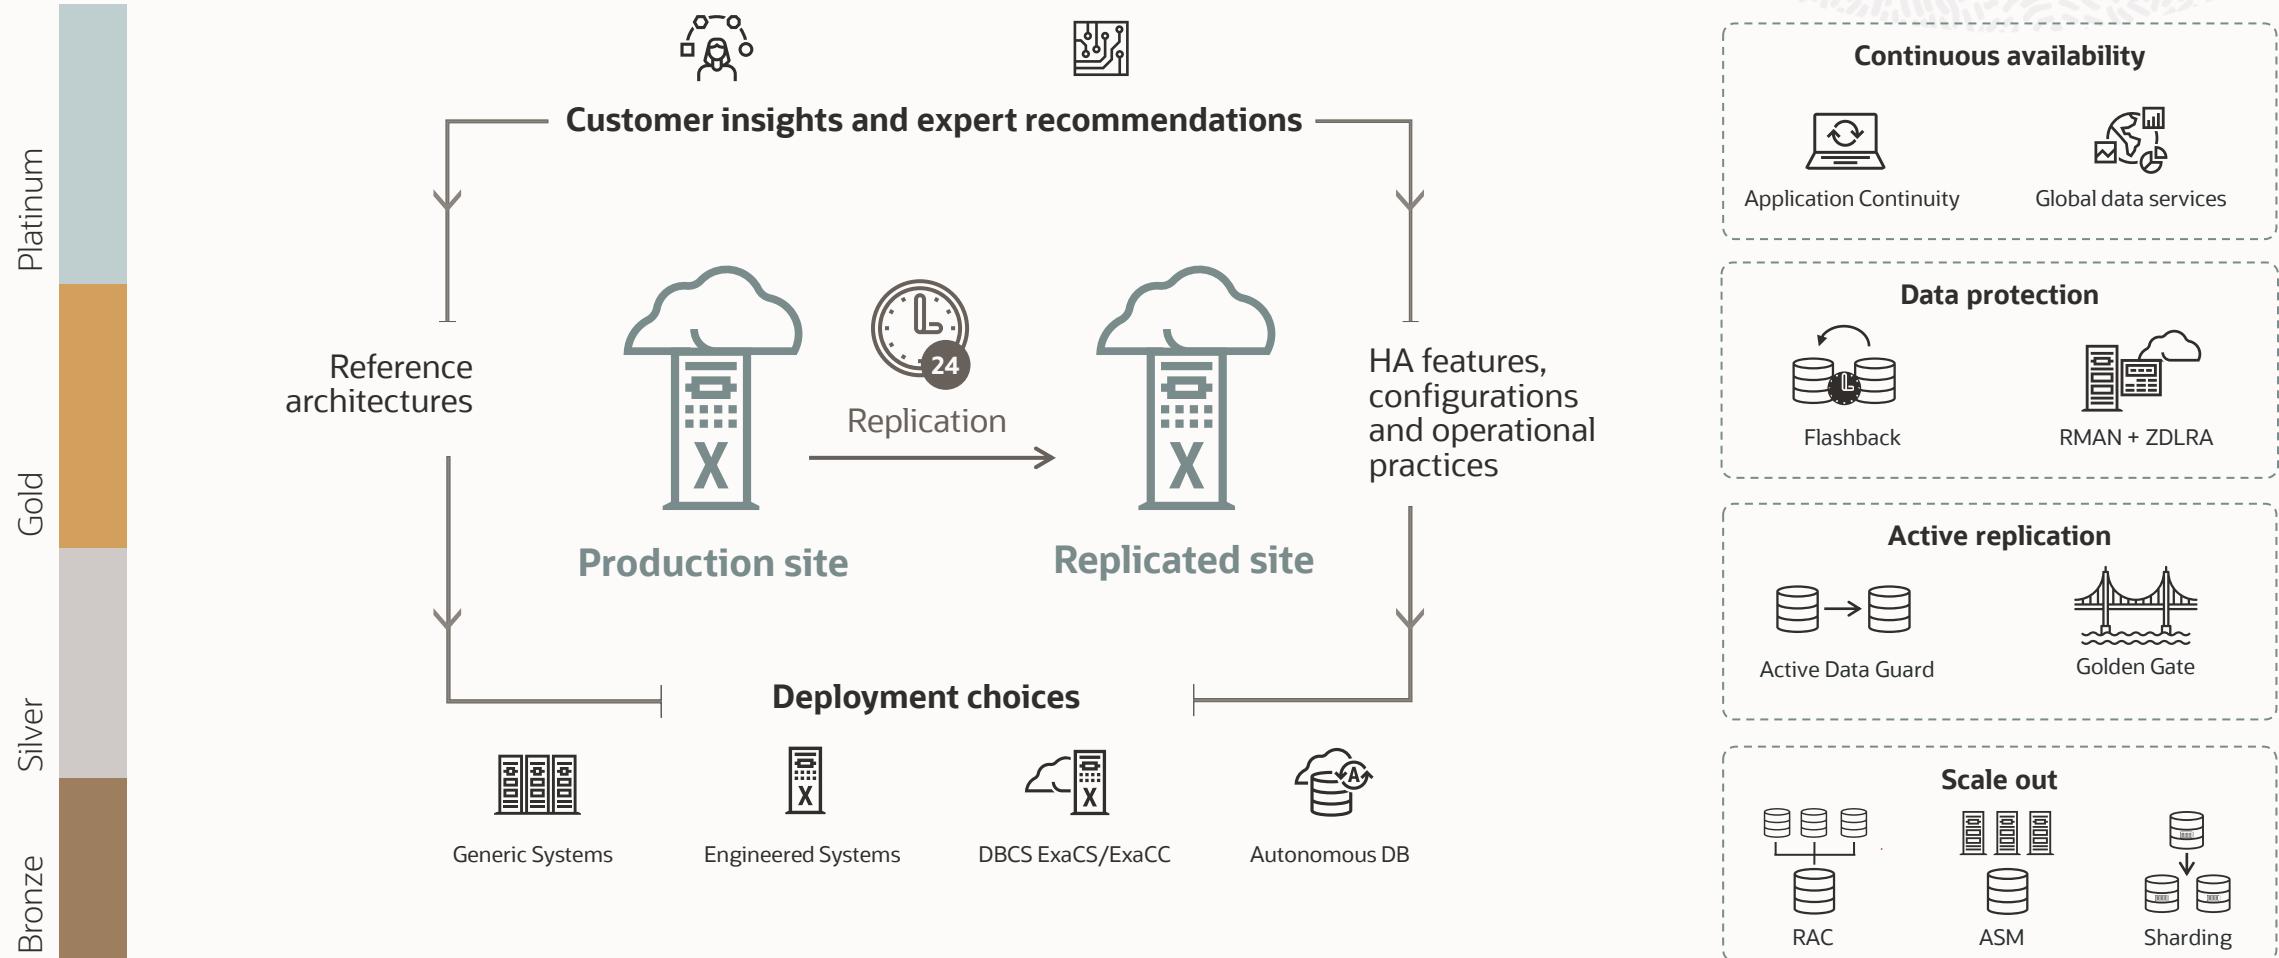
Oracle makes it simple for Apps to deliver Data Sovereignty or Massive-Scale using native Sharding

Shard monolith databases into a farm of smaller databases

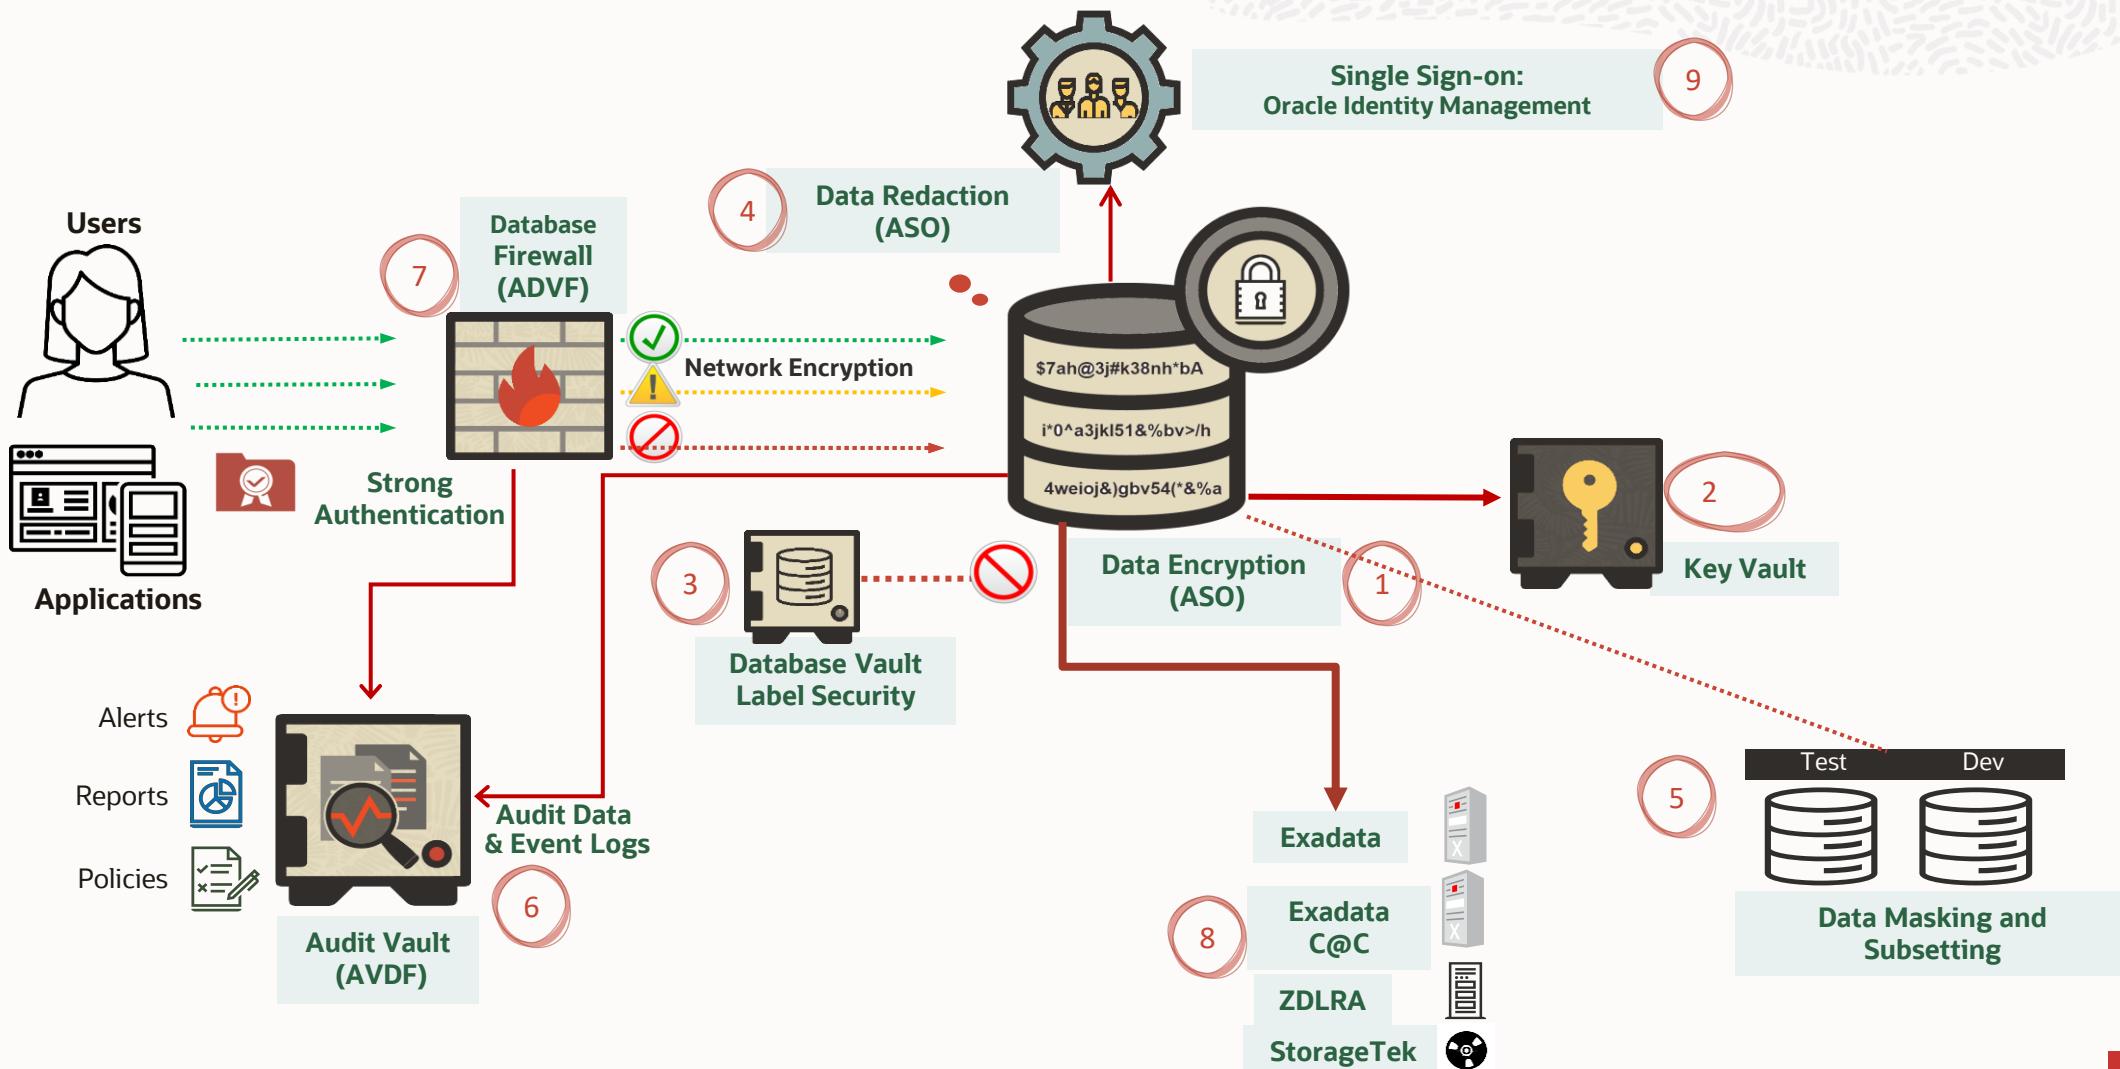
- Shards can be placed in-country to satisfy data sovereignty
- Shards are fully isolated - **linear scalability of data and users**
- Routes SQL based on shard key, or runs cross shard SQL
- Online addition and reorganization of shards

All the benefits of sharding with all the benefits of a mature SQL Database

Oracle's unique approach: Maximum Availability Architecture (MAA)



Oracle's unique approach: Maximum Security Architecture



SATLOG modernizes fleet management at scale on Oracle Cloud

“SATLOG achieves better integration and performance making fleets more productive with Oracle Autonomous Transaction Processing and its included Spatial, APEX, and JSON services on Oracle Cloud Infrastructure.”

Juergen Stausberg
CEO, SATLOG

Business Challenge:

Germany-based SATLOG helps increase the productivity of fleets across a variety of industries with various delivery requirements. Integrating data from multiple systems was a challenge for the company. It needed near real-time location and telematics data in dashboards.

Results:

Oracle Autonomous Transaction Processing helps manage data from mixed sources, such as JSON, geospatial, and relational data.

- ✓ Reduced operational costs
- ✓ Machine learning and automation is applied for provisioning, tuning, scaling, patching, and encrypting that reduces costs, time, and risk
- ✓ Using Oracle APEX and Oracle Spatial, it developed new analytical capabilities
- ✓ The company can measure, compare, and optimize costs for order processing, proof of delivery, routes, tour planning, and idle times

Products Used:

Oracle Cloud Infrastructure

Oracle Autonomous Database

APEX - Application Express

Implementation Partners: HERE Technologies; MEKRATronics; Continental Tire Pressure System CPC

Read full story [here](#)

Image courtesy of <https://www.satlog.de/en/home/>

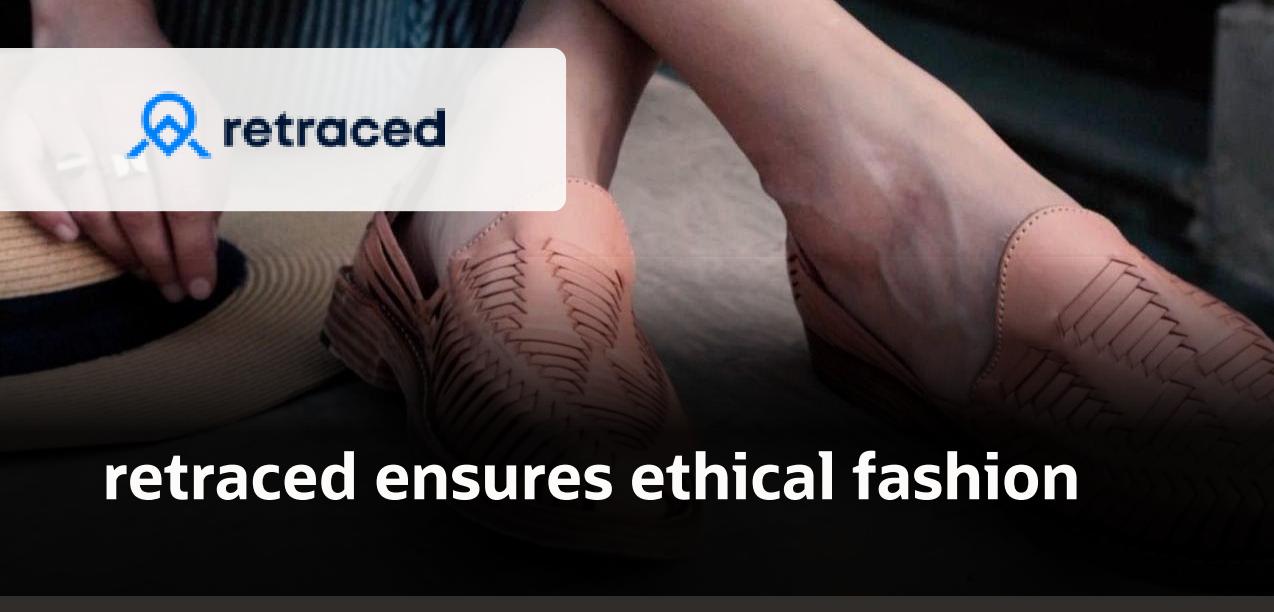


retraced ensures ethical fashion

“Having the infrastructure, database, and blockchain application running on one platform made it so much easier for us to expand our platform very quickly and at scale.”

Peter Merkert

Cofounder and CTO, retraced



Business Challenge:

German startup for sustainable sourcing of apparel, footwear, and jewelry fashion brands. It needed a low cost, easy to use platform to reassure customers that they can trust that the data in the app is reliable and verified.

Results:

- ✓ Databases autoscale capacity, backs up, **with no database admin**
- ✓ AJD and Kubernetes Engine **focuses resources more on developing business solutions**
- ✓ The solution **stores a variety of JSON documents** including copies of orders, images associated with products, components, brands, factories, farms, etc.
- ✓ Blockchain ensures data collected at every step of each customer's supply chain is reliable
- ✓ OAC provides quick insights to supply chain key performance indicators

Products Used:

Oracle Autonomous Database for transaction processing and mixed workloads (ATP)

Oracle Autonomous Database for analytics and data warehouse (ADW)

Oracle Autonomous JSON Database (AJD)

Oracle Blockchain Platform Cloud Service

Oracle Container Engine for Kubernetes

Oracle Analytics Cloud (OAC)

Read [story](#)



NEOS builds new cloud native app

“We decided to design and architect our new native SaaS app, CloudVane, on Oracle Cloud Infrastructure with Oracle Autonomous Transaction Processing at its core from the get-go for fast time to market, low administration and costs, and high performance. Now, we can add more value for our customers with resource automation, as well as data and machine learning predictions that monitor and optimize their costs.”

Davorin Capan
CEO, NEOS



Business Challenge:

Managed services provider, NEOS, wanted to create CloudVane SaaS app from scratch to add more business value with clients. It required a platform that could collect and process multiple types of data very large amounts for simpler and faster deployment than an on-premises one.

Results:

- ✓ Built new CloudVane app on ATP with Kubernetes Engine for faster time from development to deployment
- ✓ ATP enabled **integration of multiple types of data in very large amounts, (e.g., traditional structured transactions, unstructured data in documents, JSON files, etc.)** and enabled simpler and faster deployment
- ✓ Easy to provision new environment in minutes compared to hours or days on-prem and with high performance
- ✓ Database administration dropped from **80 hours per week to 4 eliminating all maintenance**
- ✓ **76% cost savings for dev and test plus 30% cost savings for production** from autoscaling and optimized storage
- ✓ In-database Machine Learning for more precise and 40%-60% faster predictions and recommendations

Products Used:

Oracle Autonomous Database for transaction processing and mixed workloads (ATP)
Oracle Machine Learning
Oracle Container Engine for Kubernetes

Read [story](#)

Summary



Modern applications need to generate value from data in new ways

- They are built using new development methodologies and technologies
- But it complicates Database Architecture by using multiple single-purpose databases
- Development focused on Integration

Oracle Database allows for easy start developing document-oriented applications

- Provides multiple APIs
- Fully supports document-oriented data model based on JSON specification
- Provides MongoDB API for MongoDB applications
- Provides multiple enterprise-scale solutions, which increase its availability, security and performance

Oracle Database vs MongoDB limitations



Limitation	Oracle Database (all editions)	MongoDB (all editions)
Max Document Size	32 MB	16 MB
Nested Depth for Documents	1024 levels	100 levels
Indexes per collection	unlimited	64
Compound index fields	unlimited (with JSON SEARCH INDEX)	32
Full document index	JSON SEARCH Index	x
Server-side functions	Functions, Procedures, Triggers	Not recommended as per MongoDB doc
Multi-document transactions	Always ACID	ACID only upon request via explicit API calls
Transaction duration	unlimited	60 seconds default
Transaction size	unlimited	<=1000 documents recommended
Aggregation data size	unlimited	100 MB RAM + explicit allowDiskUse param

Some Developer Resources for free

- [Oracle Database Express Edition \(Version 21c\)](#)
- [Pre-Built Oracle DB Developer VM](#)
- [Docker Support](#)
- Live SQL: without installation learn and share <http://livesql.oracle.com>)
- [Autonomous Database Free Trial](#)
- Free SQL courses:
[Databases for Developers: Foundations](#)

