



# Oracle Converged Database

Developing document-oriented applications  
with JSON, REST and SODA

**Witold Świerzy**

EMEA Core/Converged Database Expert

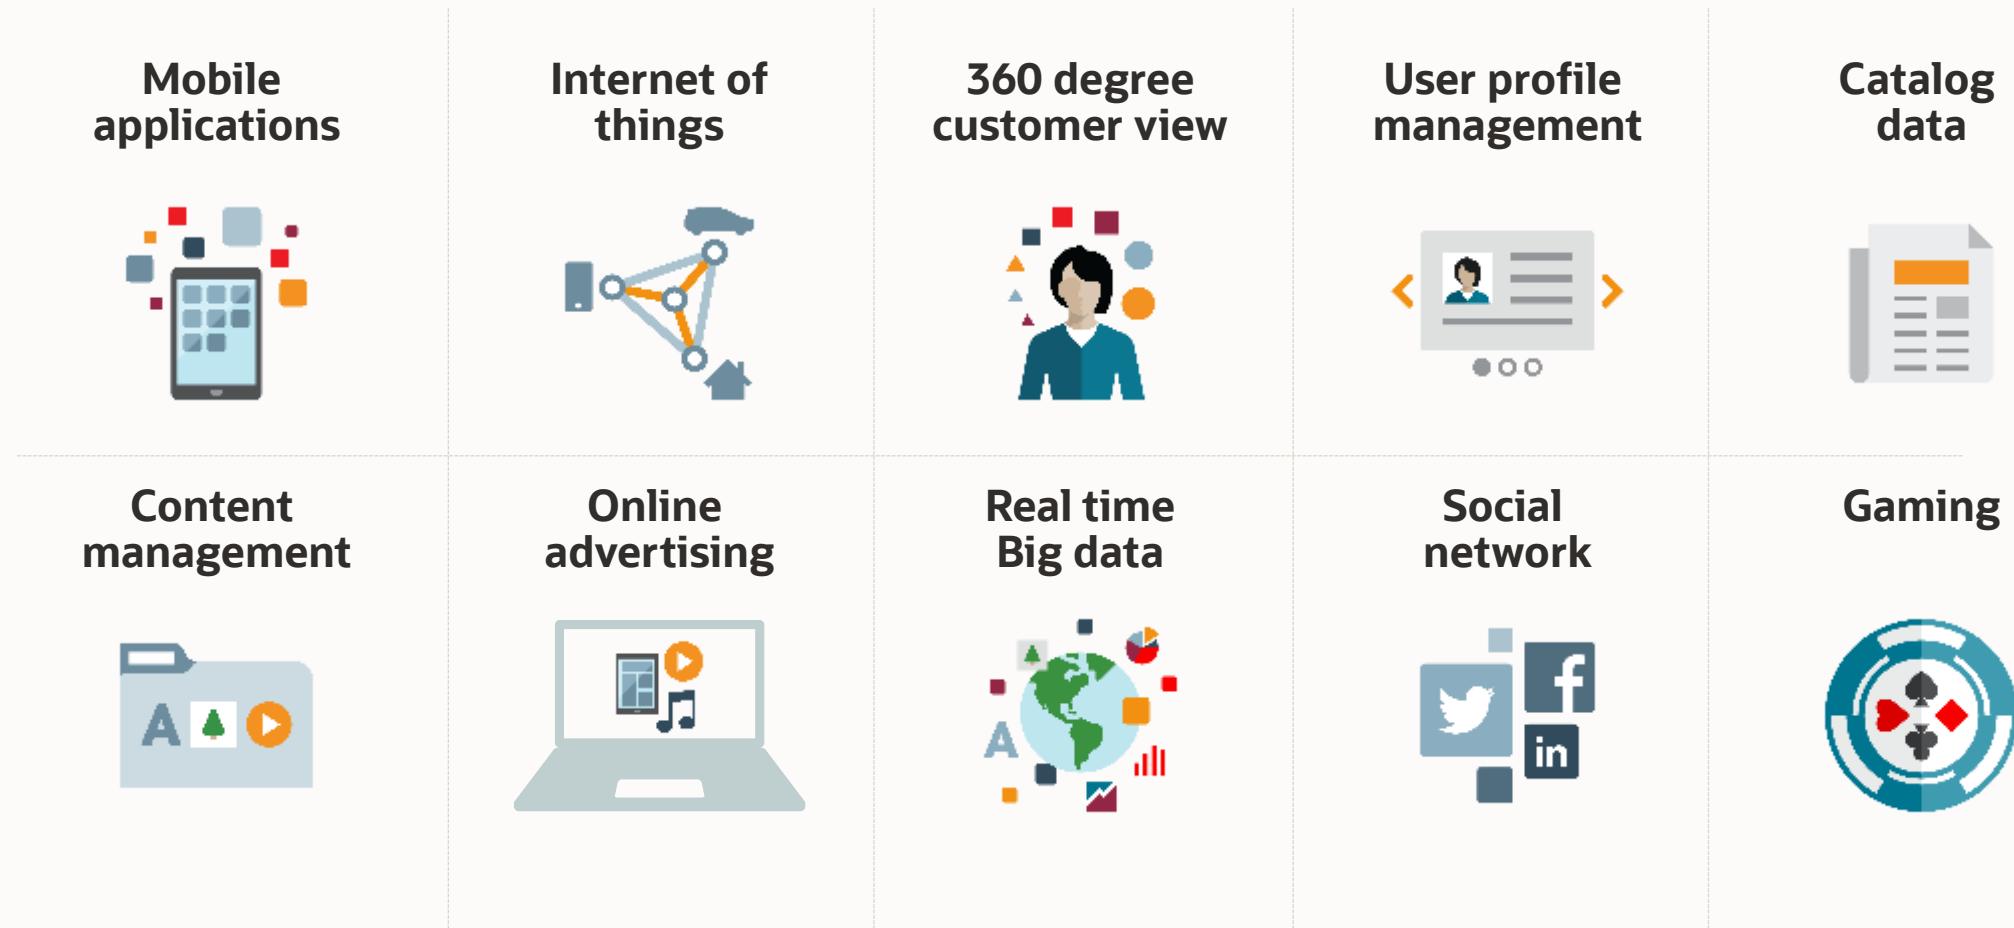
# Agenda

- **Introduction**
- **Oracle Converged Database – document oriented database**
- **MongoDB API Live Demo**
- **Summary**

# Introduction

---

# Modern Applications need flexible schema



# Why JSON can be an answer ?

- JSON stands for JavaScript Object Notation
  - Defined by standards ECMA-404 and ECMA-262
- JSON documents have
  - Fields
  - Values (Strings, Numbers, Booleans, ISO 8601 Dates, Nulls, Objects, Arrays)
- JSON documents are stored in Collections
  - Each document may have **its own Schema**
  - Schemas are flexible



```
{  
    "MyPONumber": 1602,  
    "PrettyRequestor": "Clark Kent",  
    "RequestedAt": "2020-07-22T05:18:48Z",  
    "ShippingInstructions": {  
        "Address": {  
            "PurchaseOrderNumber": 1601,  
            "KindRequestor": "Peter Parker",  
            "RequestedAt": "2020-07-21T15:26:37Z",  
            "ShippingInstructions": {  
                "Address": {  
                    "PONumber": 1600,  
                    "Requestor": "Alexis Bull",  
                    "RequestedAt": "2020-07-20T10:16:52Z",  
                    "ShippingInstructions": {  
                        "Address": {  
                            "street": "200 Sequoia Avenue",  
                            "city": "South San Francisco",  
                            "zipCode": 94080,  
                            "geometry": {  
                                "type": "Point",  
                                "coordinates": [ 37.6621  
                            ]  
                        },  
                        "Phone": [  
                            {  
                                "type": "Mobile",  
                                "number": "415-555-1234"  
                            }  
                        ],  
                        "Special Instructions": null,  
                        "AllowPartialShipment": false,  
                        "Items": [  
                            {  
                                "Description": "One Magic Christmas",  
                                "UnitPrice": 19.95,  
                                "UPCCode": 13131092899,  
                                "Quantity": 1.0  
                            }  
                        ]  
                    }  
                }  
            }  
        }  
    }  
}
```

# Why JSON can be an answer ?

## 1: JSON easily maps to Business Objects

- no need to map object to multiple tables
- single insert/query/update
- less complex SQL



### Customer: Major US Insurance:

- every insurance policy is a JSON doc
- different policy shapes
- all policy information in one doc
- no decomposition/joins
- low latency puts/gets
- great for web-application
- Relational views
  - reporting
  - daily risk analysis



## 2: JSON is schema flexible

- adding additional fields
  - *add a middle name or birth name*
- cardinality change
  - *1 address -> multiple addresses*
- Productivity, Agility

```
{  
  "firstName" : "John",  
  "middleName": "Jeremy",  
  "lastName"  : "Smith",  
  "age":25,  
  "address": [ { ...}, {...} ]}
```

### Customer: Major Global Stock Exchange

- metadata for financial products
- 1600 possible attributes
- 100 used on average
- application adds new attributes if needed without database op
- attributes in JSON column
- Custom reporting for banks
- SQL/JSON operations



## 3: JSON Document Store APIs

- Avoid SQL for simple (CRUD) operations
  - insert/find/update/delete a person
  - noSQL-style document APIs

```
SODA get -f {"lastName": "Smith"}
```

versus

```
SELECT id, data  
FROM persons  
WHERE data.lastName = 'Smith'
```

- Preserve SQL for reporting, analytics, ML

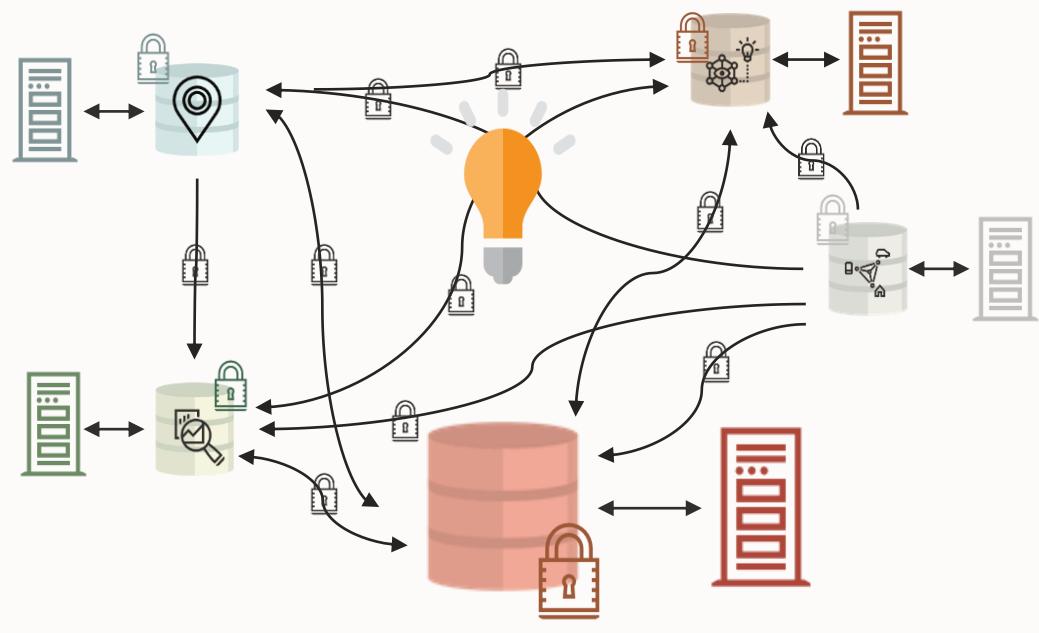
### Customer: Global Retailer, stores and online shop

- Point of sale application
  - order status tracking
  - payment tracking
  - gift card processing
- Developers preferred document store model
  - agile, productive, easy to use
- DBA and Ops preferred Oracle
  - transactions
  - SQL for reports
  - security features,...

NoSQL

SQL

# Multiple single-purpose databases may drive to complex and unmanageable architecture



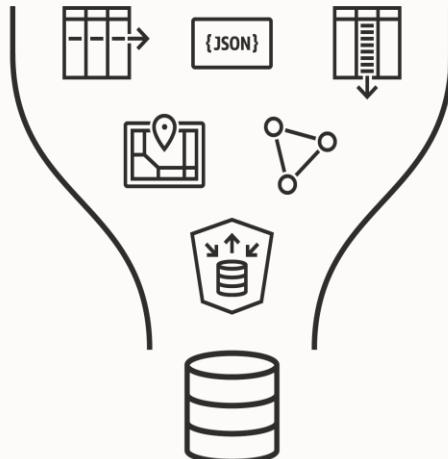
→ ETL, replication, events etc.

**Traditional paradigm assumes using many single-purpose databases drives to**

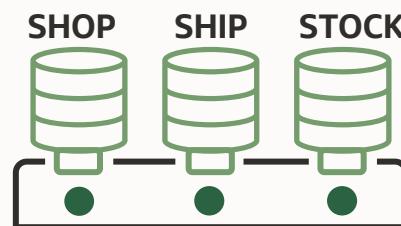
- High level of complexity of maintenance and development process
- High TCO Costs

# Oracle Database

A Converged, Open SQL Database



**Converged Database**



**Support for microservices**

## Multi-model

Best-of-Breed Relational, JSON, Spatial, Graph, Cube, Text, Blockchain  
Cross-model operations enables you to easily create value across all your data

## Multi-workload

High Performance Transactions, DW, Analytics, ML, IoT, Streaming, Multitenant  
Deep optimizations deliver exceptional price-performance across all workloads

## Most productive for developers and analysts

Same SQL and transactions operate on any data and workload  
Integrated microservices, events, REST, CI/CD, Low-code

## Support for microservices

Low-level data types and workloads should not dictate your architecture

# Oracle Converged Database Document-oriented database

---



# Oracle Converged Database

support for all modern Languages/Drivers/Tools/APIs

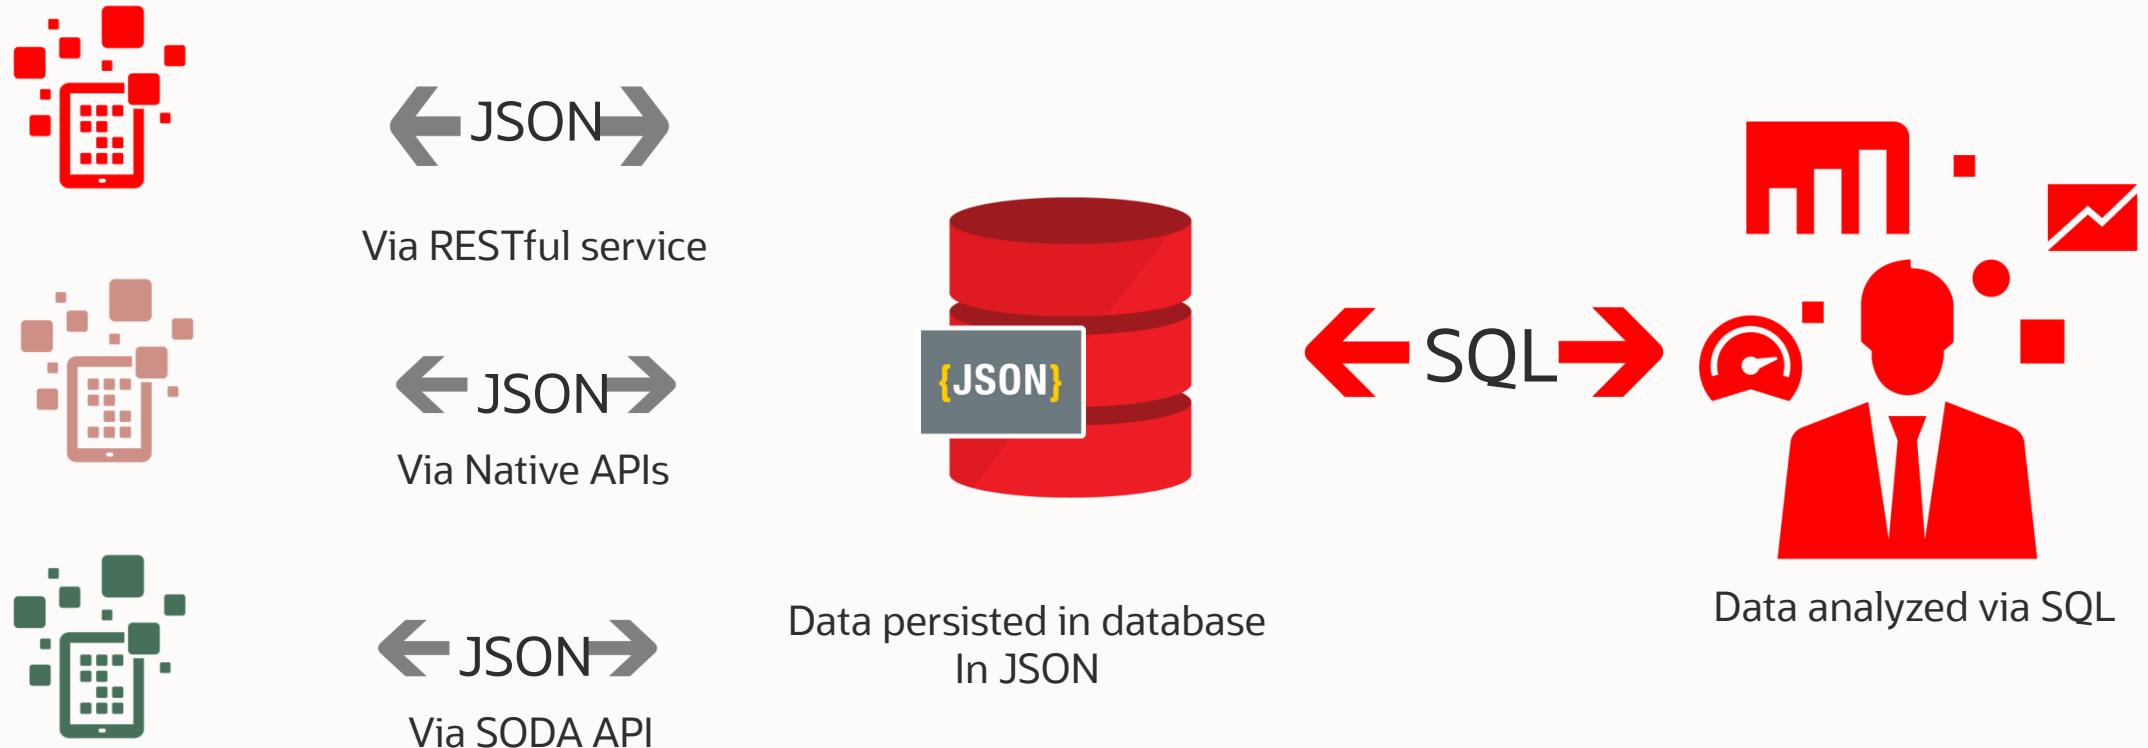


C  
C++  
Java  
.Net  
JavaScript  
Python  
PHP  
R  
Go  
Rust  
Ruby  
Perl

- SQL and PL/SQL
  - SQL Developer, Data Modeler
  - VS Code plugin
  - SQLcl (modern sqlplus)
    - With *Advanced Liquibase*
- Collections API
- SODA drivers
  - JSON results
- Data As a Microservice with REST API
- JSON Data through HTTPS

# Oracle Converged Database

Access JSON data via multiple API(s)



# Oracle Converged Database

- Full SQL support
- ACID transactions
- JSON search indexes
- Advanced security
- APEX low-code development



# Oracle Converged Database



- You have the choice to store JSON data as
  - Collections
  - Column(s) in relational tables
- JSON documents are stored as
  - Native Binary data type
    - **JSON** (up to 32 MB)
  - Textual data type
    - VARCHAR2 (up to 32\* KB)
    - BLOB (up to 2 GB)
    - CLOB (up to 1 GB)



## Support for MongoDB

### Develop and run MongoDB workloads in the Oracle Autonomous JSON Database



#### Modern document-centric development

- JSON Collections-based data model
- Rich clients – MongoDB API, REST and SODA based development API
- Native JSON storage with advanced indexes and optimized performance

#### ... and proven enterprise functionality

- ACID Transactions
- SQL-based Reporting and Analytics (including scalable parallel execution)

#### ... running on Oracle Database platform

- Availability
- Security
- Elasticity

# Support for MongoDB

Develop and run MongoDB workloads in the Oracle Autonomous JSON Database



## API compatibility details

### MongoDB API version 4.2 compatible

- Transactions supported
- Load balanced connections supported like on MongoDB Atlas

### MongoDB API backed by Oracle user and roles management instead of MongoDB's

- Allows Oracle's enterprise-class security features to be used with MongoDB collections
- Makes unified user management easy

# Oracle Autonomous Database - support for MongoDB

## Develop and run MongoDB workloads in the Oracle Autonomous JSON Database

### JSON Development IDE

The screenshot shows the Oracle Database Actions JSON interface. At the top, there's a search bar with the query `{"brand": "Apple"}`. Below it, two documents are displayed in a list view:

```
{ "price": 89.99,
  "rating": 8,
  "countInStock": 3,
  "description": "Bluetooth technology lets you connect it with compatible devices wirelessly. High-quality AAC audio offers immersive listening experience. Built-in microphone allows you to take calls while working",
  "category": "Electronics",
  "createdAt": "2022-01-14T22:30:27.528000Z",
  "_id": "61e1f983746641a5d4af602a",
  "updatedAt": "2022-01-14T22:30:27.528000Z",
  "_v": 0,
  "reviews": [],
  "name": "Airpods Wireless Bluetooth Headphones",
  "numReviews": 0,
  "user": "61e1f983746641a5d4af6027",
  "brand": "Apple",
  "image": "/images/airpods.jpg"
}

{ "price": 599.99,
  "rating": 10,
  "countInStock": 10,
  "description": "Introducing the iPhone 11 Pro. A transformative triple-camera system that adds tons of capability without complexity. An unprecedented leap in battery life",
  "category": "Electronics",
  "createdAt": "2022-01-14T22:30:27.528000Z",
  "_id": "61e1f983746641a5d4af602b",
  "updatedAt": "2022-01-14T22:30:27.528000Z",
  "_v": 0,
  "reviews": [],
  "name": "iPhone 11 Pro 256GB Memory",
  "numReviews": 0
}
```

At the bottom, a status bar indicates `2 0 0 0 | 10:33:19 PM - REST call resolved successfully.`

Oracle Cloud tools

The screenshot shows the MongoDB Compass interface for the `merndemo.products` collection. It displays 7 documents with the following data:

```
price: 89.99
countInStock: 3
name: "Airpods Wireless Bluetooth Headphones"
image: "/images/airpods.jpg"
description: "Bluetooth technology lets you connect it with compatible devices wirelessly. High-quality AAC audio offers immersive listening experience. Built-in microphone allows you to take calls while working"
brand: "Apple"
category: "Electronics"
user: ObjectId("61e1f983746641a5d4af6027")
> reviews: Array
  _v: 0
  createdAt: 2022-01-14T22:30:27.528+00:00
  updatedAt: 2022-01-14T22:30:27.528+00:00

_id: ObjectId("61e1f983746641a5d4af602b")
rating: 0
numReviews: 0
price: 599.99
countInStock: 10
name: "iPhone 11 Pro 256GB Memory"
image: "/images/phone.jpg"
description: "Introducing the iPhone 11 Pro. A transformative triple-camera system that adds tons of capability without complexity. An unprecedented leap in battery life"
brand: "Apple"
```

MongoDB Compass and other tools

# Support for MongoDB

**Recommended tools and driver versions or higher with support for load-balanced connections**



- **C 1.19.0**
- **C# 2.13.0**
- **Compass 1.28.1**
- **Database Tools 100.5.0 (includes mongoexport, mongorestore, and mongodump)**
- **Go 1.6.0**
- **Java 4.3.0**
- **MongoSH 0.15.6**
- **Node.js driver 4.1.0**
- **PyMongo 3.12.0 (for Python language)**
- **Ruby 2.16.0**
- **Rust 2.1.0**

Source: <https://docs.oracle.com/en/database/oracle/mongodb-api/mgapi/overview-oracle-database-api-mongodb.html#GUID-9B704340-5E1D-4CA9-8944-6296A69221B0>

# Oracle Converged Database

## SQL support for JSON

- **SQL** for advanced reporting and analytical operations on JSON documents
  - Dot notation
    - SELECT coll.doc.address.postalCode...
  - SQL/JSON Path Expressions
  - Rich set of functions
    - Views + On-Demand Refreshable Materialized Views
    - Analytical functions (over (), window (...))
    - Documents Pattern Matching (match\_recognize)
    - JSON documents generation from relational data
    - ...
- **PL/SQL** for **server-side** processing
  - Parse and generate Documents
    - apex\_json
  - Secure access to JSON Open Data
    - apex\_web\_service

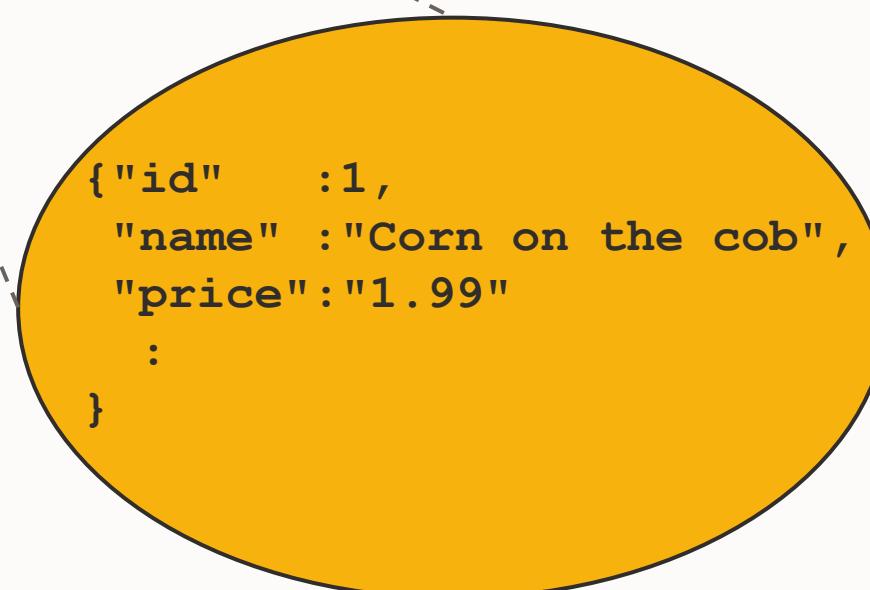


# Oracle Converged Database

## SQL support for JSON

```
-- 1. Create a menu-items table  
-- where menus items are stored as JSON
```

```
CREATE TABLE menu_items(  
    rest_name  VARCHAR2(255),  
    menu-item  JSON);
```



```
{"id" : 1,  
 "name" :"Corn on the cob",  
 "price": "1.99"  
 :  
 }
```

-- 2. Use simple dot notation to access elements within the JSON docs

```
SELECT m.menu_item.name item,  
       m.menu_item.price price  
  FROM menu_items m;
```

ITEM	PRICE
-----	-----
Corn on the cob	1.99

-- 3. Create index on a JSON column element to speed up rows filtering

```
CREATE INDEX menu_items_idx ON  
purchase_orders  
( JSON_VALUE(JSON_DOCUMENT, '$.id.number()' ));
```



# Oracle Converged Database

## SQL support for JSON



### Native Support to Index and Search JSON Documents

Oracle makes it simple for Apps to index, search and analyze text

Search text using keyword search, context queries, pattern matching, etc.

- Search websites, catalogs, documents, LOBs

Perform linguistic analysis on documents to easily classify them

- Classify customer feedback as positive, negative, or neutral using sentiment analysis

Search JSON or XML documents using the structure of the document to restrict the search

- Find all purchase orders where the comment field contains “damage” and “delivered”

# Oracle Converged Database

## SQL support for JSON



### Native Support to Index and Search JSON Documents

-- 1. Create a text index on the JSON column of the purchase\_orders table

```
CREATE SEARCH INDEX PO_search_ind ON purchase_orders(po_doc) FOR JSON;
```

--2. Search for PO documents with “Delivered” and “Damage” in the comments field

```
SELECT po_number, po_doc.comment comment
FROM   purchase_orders
WHERE  JSON_TEXTCONTAINS(po_doc, '$.comment', 'Delivered and Damage');
```

# Oracle Converged Database

## SQL support for JSON

### Ensuring consistency

- Constraints

```
CREATE UNIQUE INDEX unique_idx ON purchase_orders
( JSON_VALUE(JSON_DOCUMENT, '$.PONumber.number()' ));  
  
SQL> soda insert purchase_orders {"PONumber": 1601 };  
SQL> commit;  
  
SQL> soda insert purchase_orders {"PONumber": 1601 };
```

**ORA-00001: unique constraint (HR.UNIQUE\_IDX) violated**

```
ALTER TABLE purchase_orders ADD CONSTRAINT
po_address_zipcode_positive_number CHECK (
    JSON_VALUE( json_document,
        '$.ShippingInstructions.Address.zipCode'
        returning number)
    > 0 )
ENABLE      -- enable constraint
NOVALIDATE; -- but don't check existing documents
```

- Triggers

```
CREATE TRIGGER validate_items
AFTER UPDATE ON purchase_orders FOR EACH ROW
DECLARE l_number_of_products number;
BEGIN
    SELECT sum(Quantity) into l_number_of_products
    FROM JSON_TABLE(:NEW.json_document, '$'
        Columns( Nested Items[*] Columns(Quantity NUMBER)));
    IF l_number_of_products > 10 THEN
        RAISE_APPLICATION_ERROR(-20001, 'No more than 10
            products ordered at the same time.');
    END IF;
END;
/  
  
SQL> UPDATE purchase_orders SET json_document =
    JSON_TRANSFORM( json_document, APPEND '$.Items' =
    '{"Description":"Shiny toy", "UnitPrice":1.0,
        "UPCCode": 123456789, "Quantity": 5}' format json );  
  
ORA-20001: No more than 10 products ordered at the same time.
```

# Oracle Converged Database

## SODA (Simple Oracle Document Access) and SQLcl



### SODA: Goals

- Enable schemaless development on top of an Oracle Database
  - Provide a simple NoSQL-style API for working with documents
- Make it easy to use Oracle as a NoSQL-style document store
  - Allow developers to work with Oracle without learning SQL
  - Allow developers to work with Oracle without DBA support
- Support all common application development environments
  - Traditional programming languages
  - Scripting languages and frameworks



# Oracle Converged Database

## SODA (Simple Oracle Document Access) and SQLcl



### SODA APIs

- NoSQL-style APIs for
  - Java, JavaScript/Node.js, Python, REST, PL/SQL, C...
- Used to manage JSON data
  - create collections
  - store documents in collections
  - retrieve documents
  - query documents
- **No need to know SQL!**

### SQLcl

- Modern SQL Developer Command Line interface for Oracle database
- Provides
  - inline editing, statement completion, command recall...
- **SODA commands**

# Oracle Converged Database

SODA (Simple Oracle Document Access) and SQLcl



## SQLcl

- Modern SQL Developer Command Line interface for Oracle Database
- Provides
  - Inline editing, statement completion, command recall...
  - **SODA commands**
  - **Full CRUD operations support**

```
SODA allows schemaless application development using the JSON data model.
SODA create <collection_name>
Create a new collection

SODA list
List all the collections

SODA get <collection_name> [-all | -f | -k | -klist] [{<key> | <k1> <k2> ... > | <qbe>}]
List documents the collection
Optional arguments:
  -all    list the keys of all docs in the collection
  -k     list docs matching the specific <key>
  -klist  list docs matching the list of keys
  -f      list docs matching the <qbe>

SODA insert <collection_name> <json_str | filename>
Insert a new document within a collection

SODA drop <collection_name>
Delete existing collection

SODA count <collection_name> [<qbe>]
Count # of docs inside collection.
Optional <qbe> returns # of matching docs

SODA replace <collection_name> <oldkey> <new_{str|doc}>
Replace one doc for another

SODA remove <collection_name> [-k | -klist | -f] {<key> | <k1> <k2> ... > | <qbe>}
Remove doc(s) from collection
Optional arguments:
  -k     remove doc in collection matching the specific <key>
  -klist remove doc in collection matching the list <key1> <key2> ... >
  -f      remove doc in collection matching <qbe>
```

# Oracle SODA

## examples

### Node.js

```
conn = await oracledb.getConnection(...);
db = conn.get SodaDatabase();
col = await db.createCollection("purchase_orders");
await col.drop();
```

### Python

```
conn = cx_Oracle.connect(...);
db = conn.get SodaDatabase();
col = db.createCollection("purchase_orders");
col.drop();
```

### Java

```
OracleClient client = new OracleRDBMSClient();
db = client.getDatabase(jdbcConn);
OracleCollection col =
db.admin.createCollection("purchase_orders");
col.admin().drop();
```

### PL/SQL (and Oracle Application Express)

```
col := dbms_soda.create_collection('purchase_orders');
select dbms_soda.drop_collection('purchase_orders')
from dual;
```

# Oracle SODA

Sample services provide by SODA for REST



GET /DBSODA/schema	List all collections in a schema
GET /DBSODA/schema/collection	Get all objects in collection
GET /DBSODA/schema/collection/id	Get specific object in collection
PUT /DBSODA/schema/collection	Create a collection if necessary
PUT /DBSODA/schema/collection/id	Update object with id
POST /DBSODA/schema/collection	Insert object into collection
POST /DBSODA/schema/coll?action=query	Find objects matching filter in body

# Oracle REST Data Services

## ORDS powers the Schema Service, A Quick Detour

- Provides data access consistent with modern App Dev frameworks
  - Mid tier application
  - Can map standard http(s) RESTful gets and posts to SQL
  - Can declaratively returns results in JSON format
  - JavaScript friendly
  - Can support high numbers of end users
- Services
  - HTTP(s) relational data access
  - Oracle JSON collection based schema-less access
  - Oracle NoSQL access over HTTP
  - Oracle APEX mid-tier, web toolkit applications, mod\_plsql replacement



# Oracle Converged Database

## Oracle REST Data Services

- Available **Standalone (Jetty)**, Weblogic, Tomcat & Glassfish
- Turns Database Service into an RESTful API service
- Fully provisioned and functional in all cloud editions
- Available since 11g, **no extra cost**
- Allows publishing of URI based access to Oracle database over REST
- Results in JSON or CSV
- Mapping of URI to SQL or PL/SQL
- All HTML methods GET, PUT, POST, DELETE, PATCH
- OAuth2 integration
- Highly scalable, can use all features of database



# Oracle REST Data Services

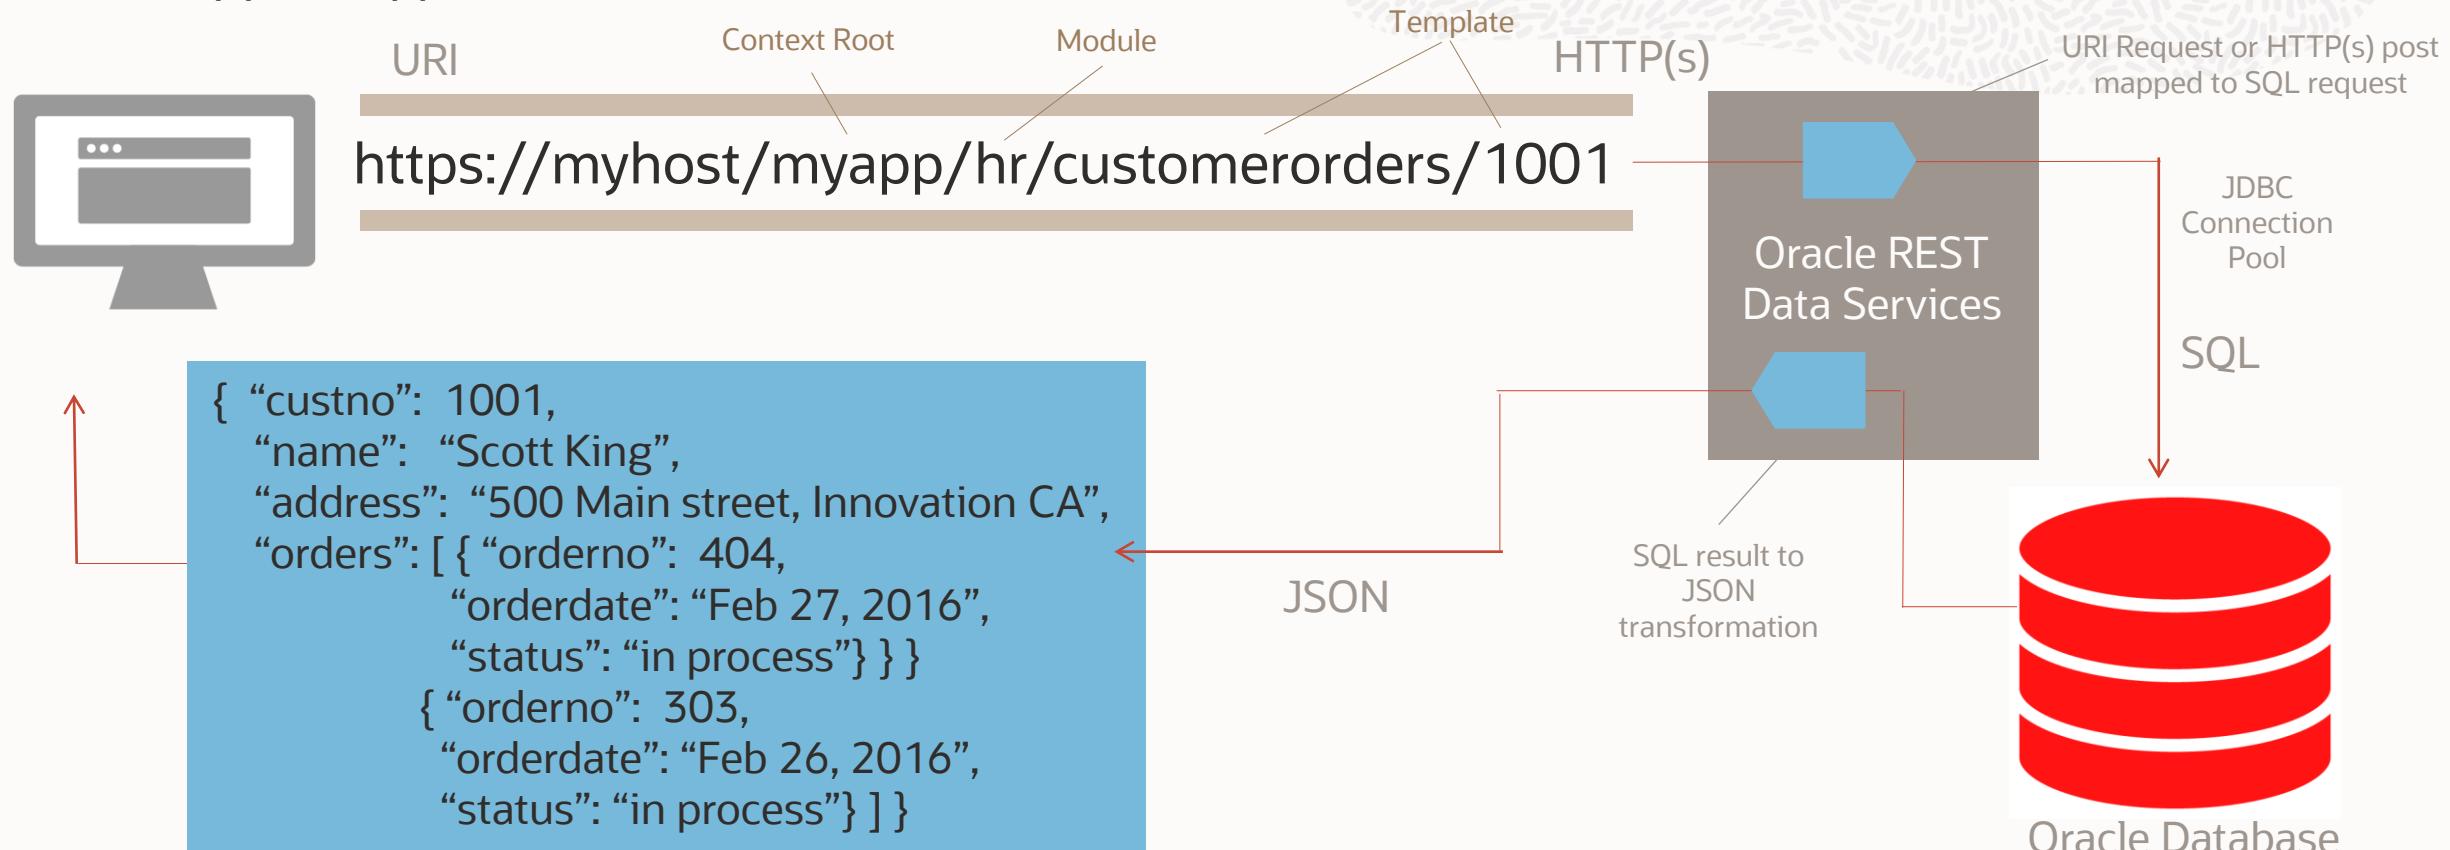
Serving JSON results from relational data



- Data stored in standard relational tables and columns
- Oracle REST Data Services (ORDS) Developer defines URI<>SQL mapping
- App Developer calls named URI over HTTP(S) gets and posts

# Oracle REST Data Services

HTTP(s) API App-Dev with Relational Tables in Oracle Database



ORDS maps standard URI requests to corresponding relational SQL (not schemaless): e.g. SQL SELECT from customers and orders table.

ORDS also transforms the SQL results into JavaScript Object Notation (JSON), other formats include HTML, binary and CSV.

Fully committed to supporting any and all standards required by Fusion / SaaS / FMW; we are actively engaged in the ongoing dialog.

# Oracle REST Data Services

RESTful development with JSON

RESTful services are a simple well, understood model

CRUD operations map to HTTP Verbs

Create/Update : PUT / POST  
Retrieve:GET – Delete :  
DELETE

Other operations, such as Query by Example, Bulk Insert and Indexing are mapped to variants of POST

JSON document forms the payload of the HTTP Request or Response

Stateless model, no transaction support

# Oracle Database

The most advanced database platform in the market

Document-oriented data model is fully supported by all the database options

## Consolidated and Mixed-Workloads



Multitenant Database, Exadata, OLTP and analytics in one

## Fault-Tolerant Availability



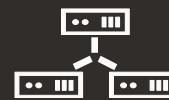
RAC, Data Guard, Flashback, Online Data Reorg, Online Patching, MAA

## Highest Security



Data Safe, Advanced Security, Encryption, Roles, Privileges ...

## Transparent Scaling



RAC, Exadata, Parallel SQL, Native Sharding

## Comprehensive Data Model Support



JSON, text, graph, spatial, blockchain

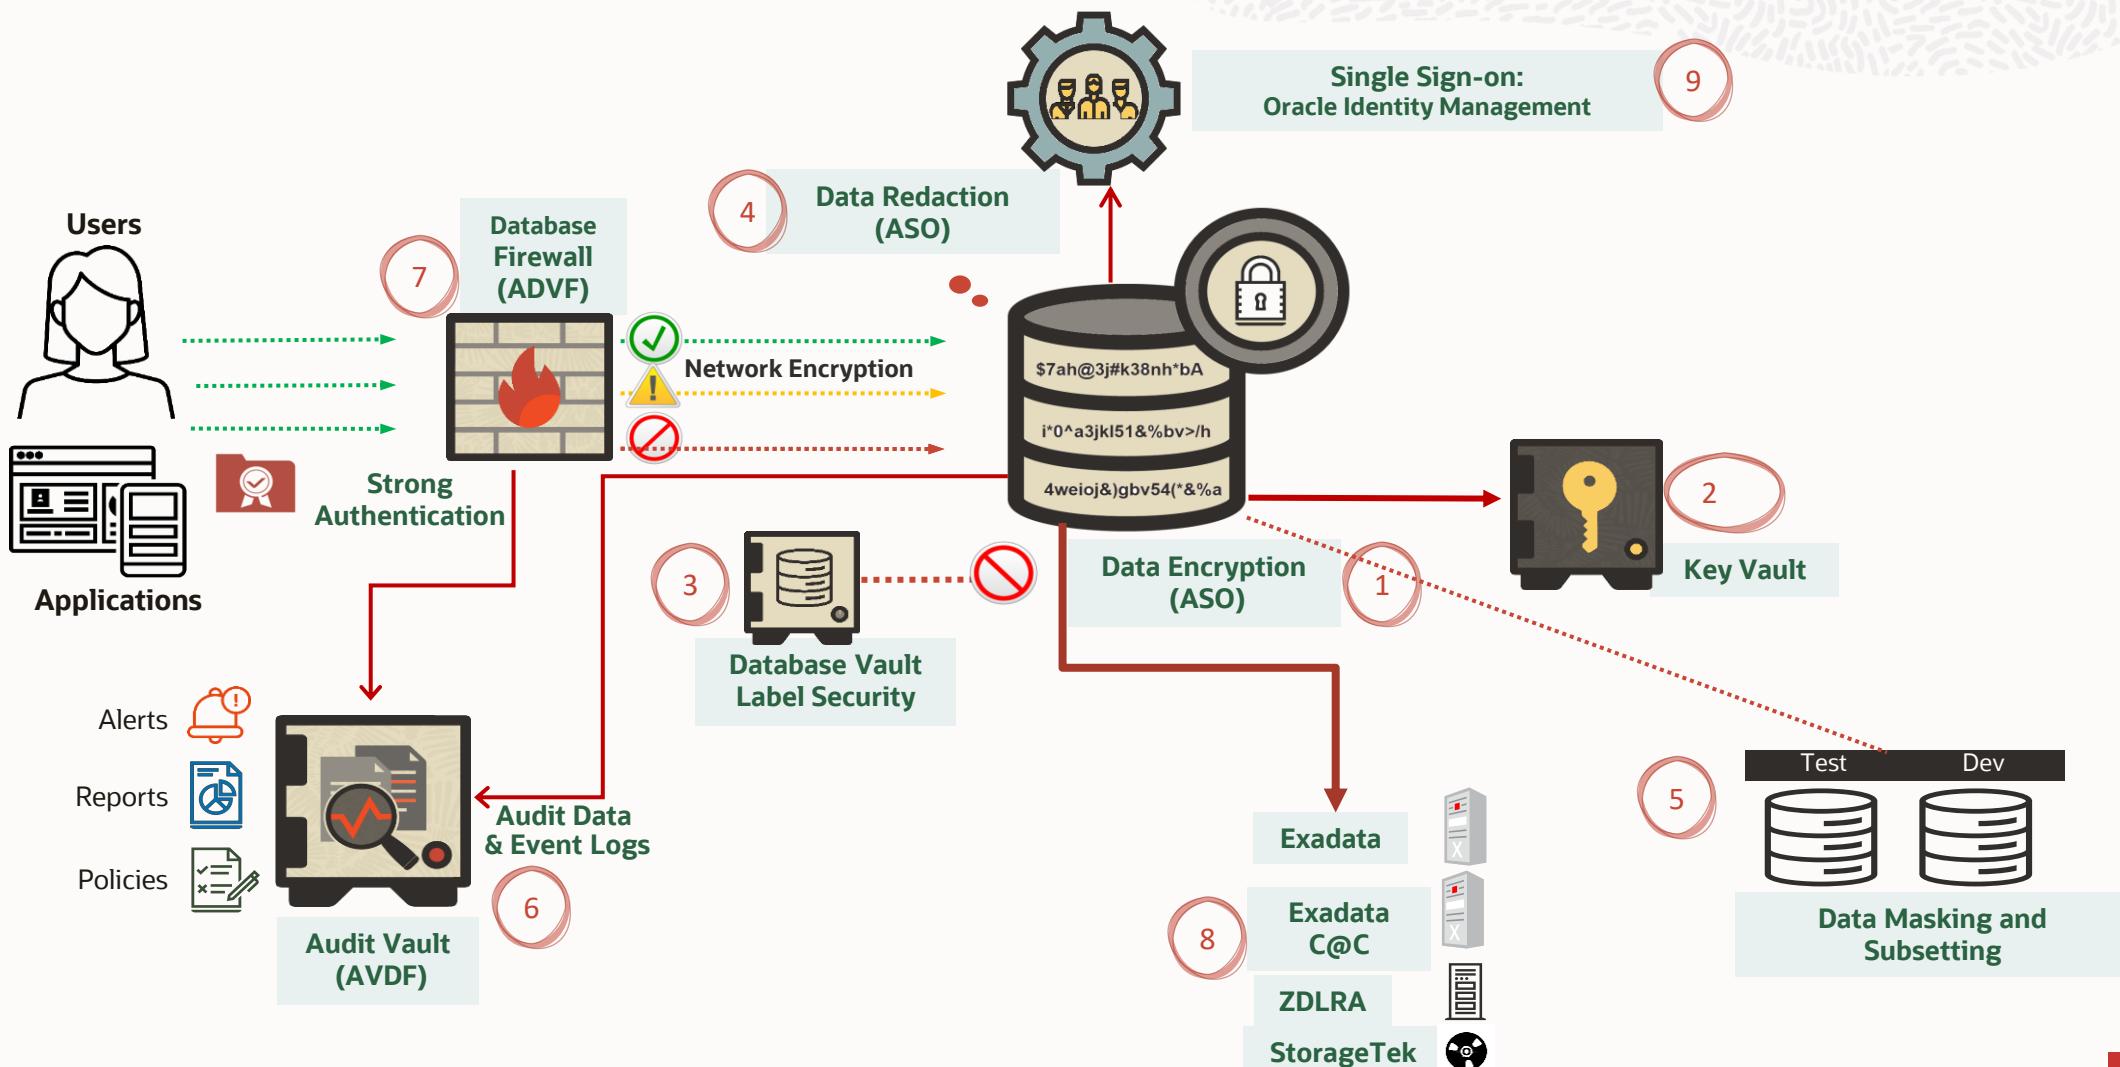
## Powerful Analytics



In-memory analytics, machine learning, data partitioning

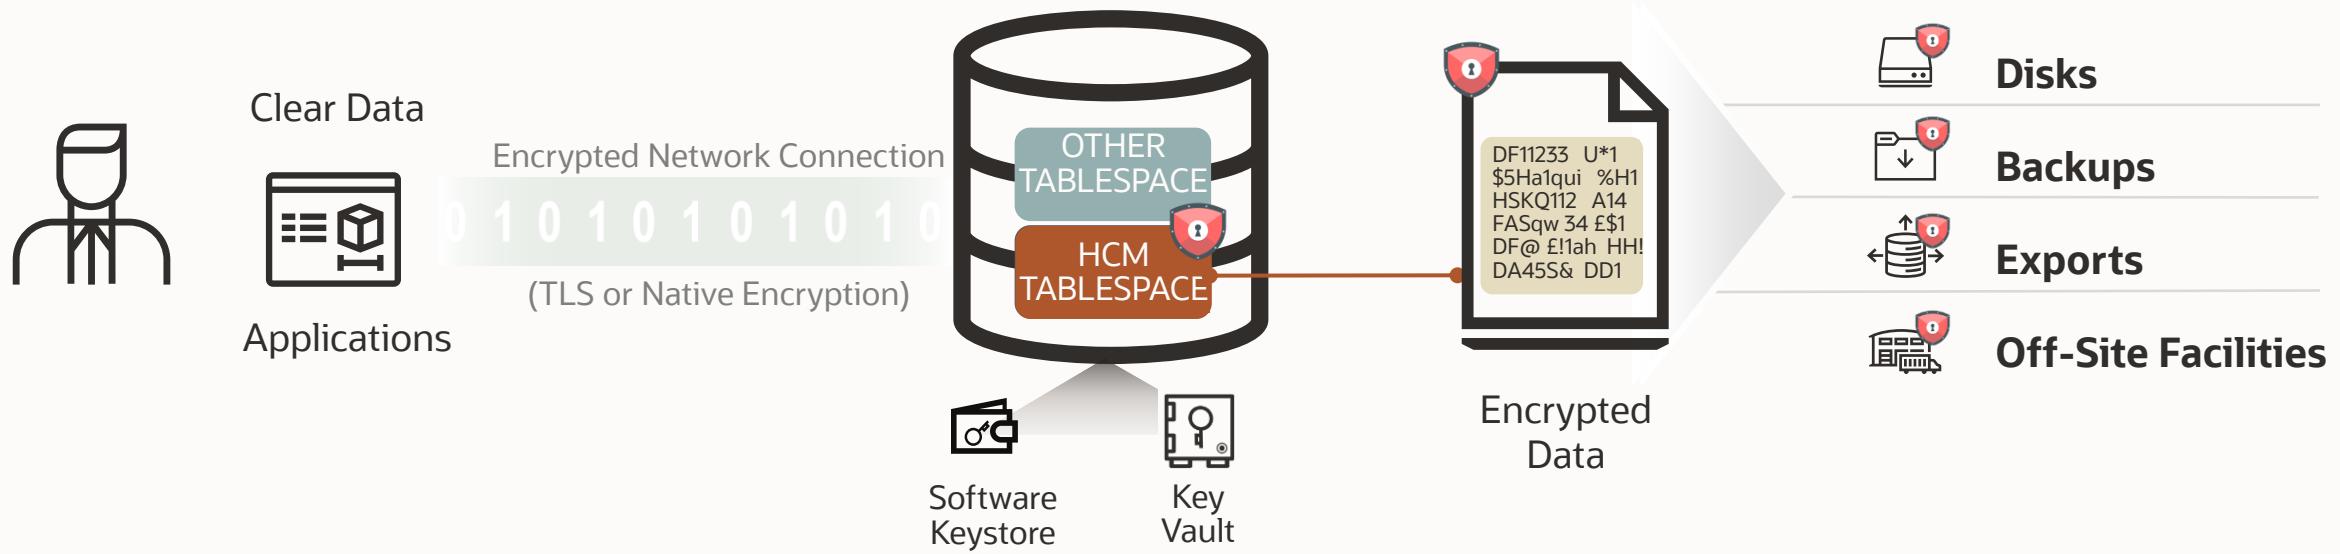
These technologies fully support document-oriented Oracle Converged Database

# Oracle's unique approach: Maximum Security Architecture



# Oracle Converged Database

## Transparent Data Encryption (TDE)



Encrypts entire application tablespaces or an application column

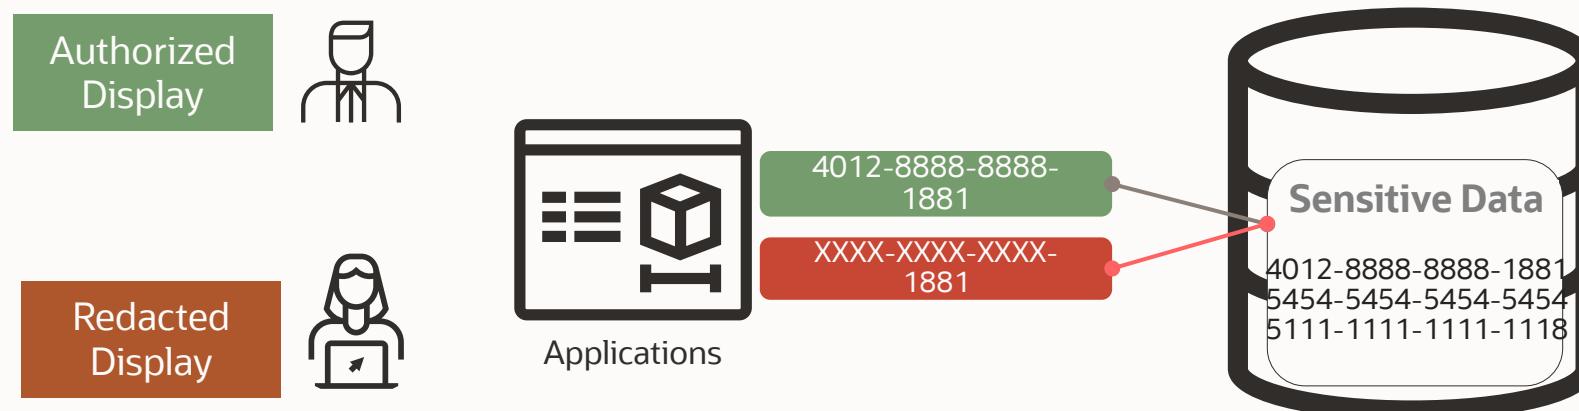
Protects the database files on disk and in backups

No application changes required

Integrated with the Oracle technology stack

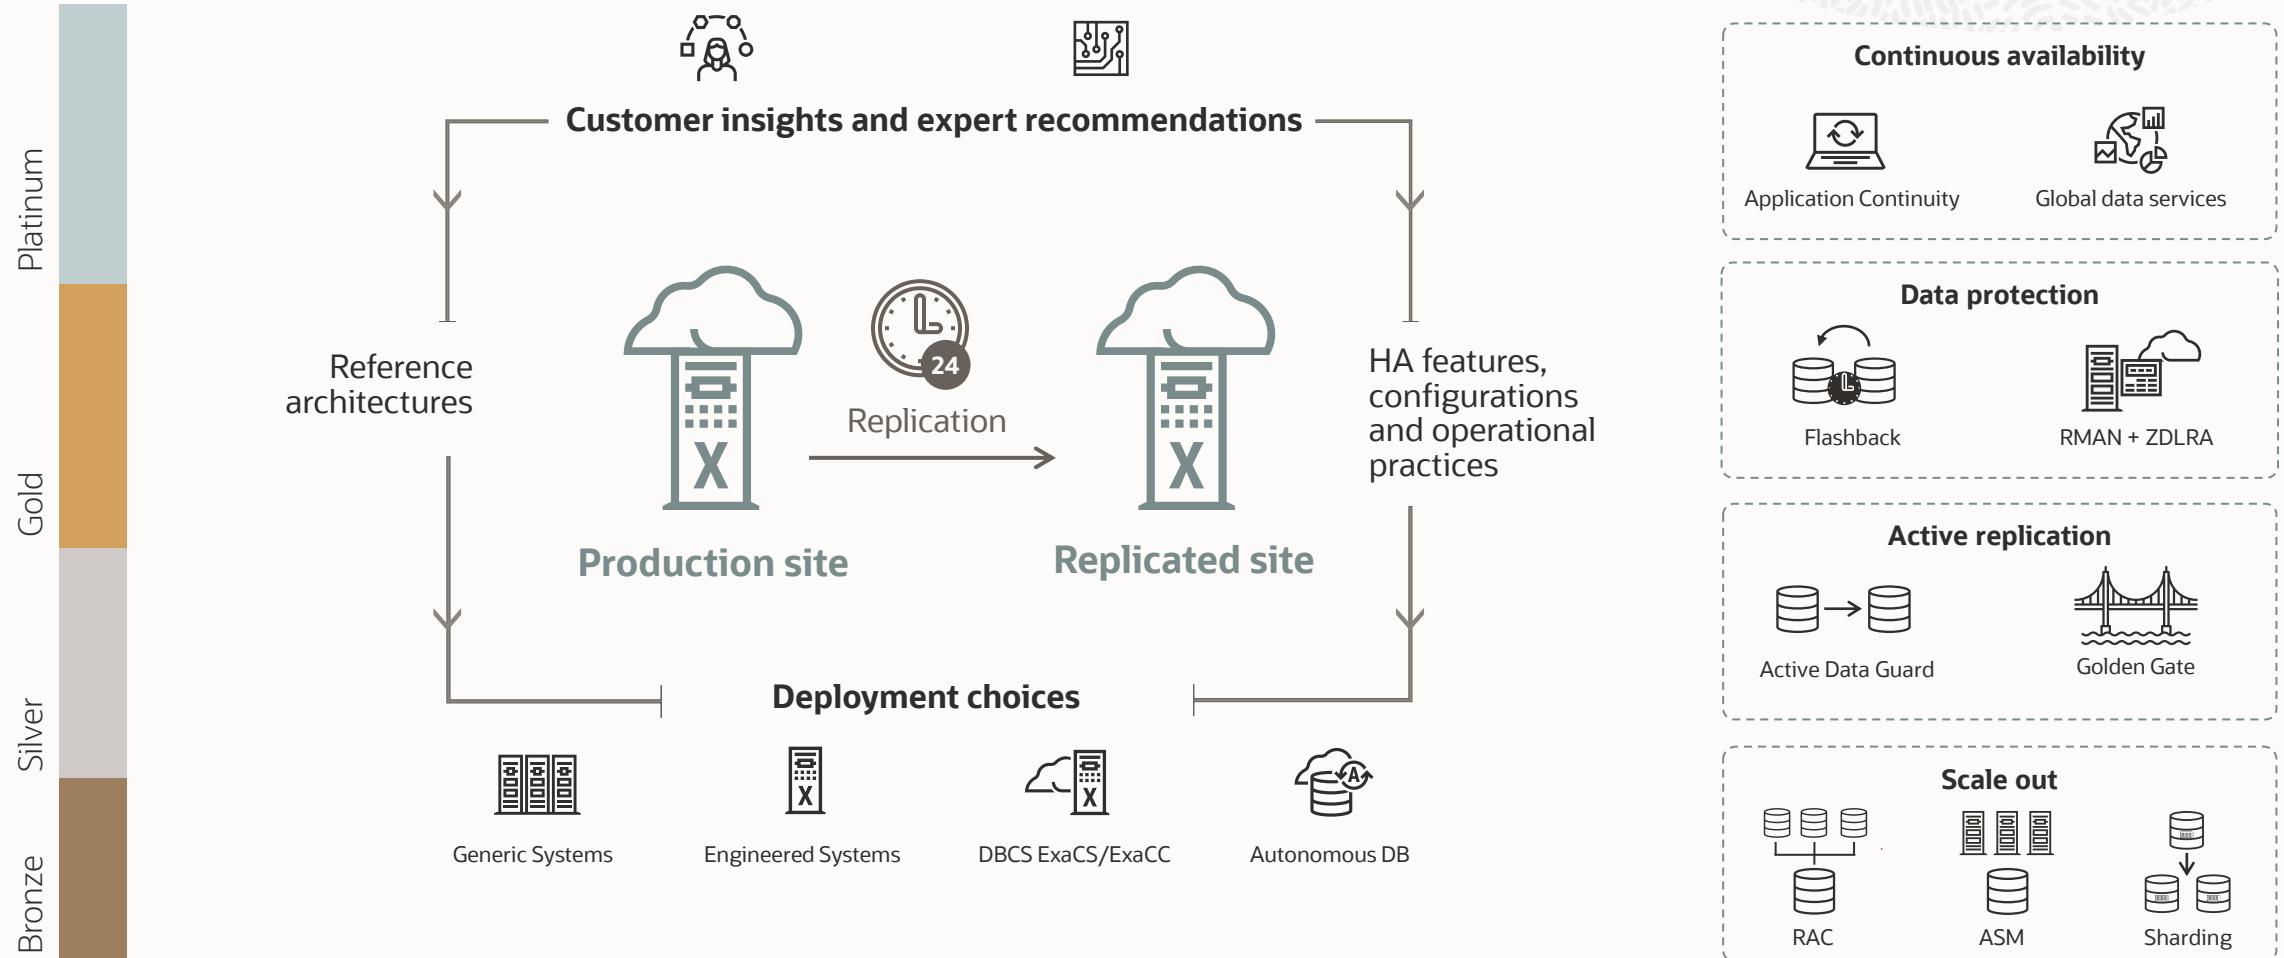
# Oracle Converged Database

## Data Redaction (Part of Advanced Security Option)



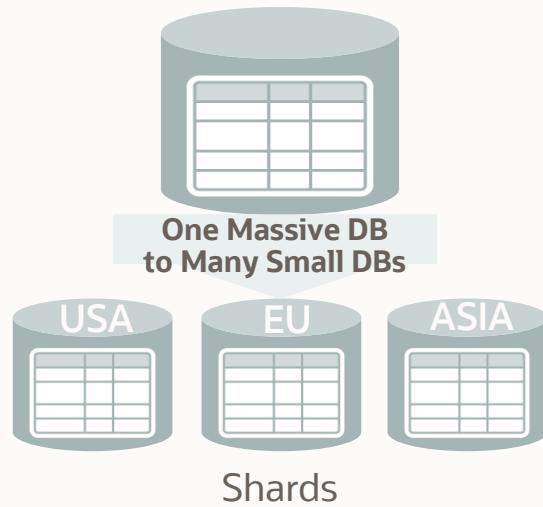
- Dynamic masking of application data based upon user name, IP, application context, and other session factors
- Library of redaction policies and point-and-click policy definition
- No impact on operational activities

# Oracle Maximum Availability Architecture (MAA)



# Oracle Converged Database

## Native Sharding Simplifies Distributed Data Architecture



Oracle makes it simple for Apps to deliver Data Sovereignty or Massive-Scale using native Sharding

Shard monolith databases into a farm of smaller databases

- Shards can be placed in-country to satisfy data sovereignty
- Shards are fully isolated - **linear scalability of data and users**
- Routes SQL based on shard key, or runs cross shard SQL
- Online addition and reorganization of shards

All the benefits of sharding with all the benefits of a mature SQL Database

# MongoDB API Live Demo

---

Live Demo of MongoDB API 23c is fully described in  
**Oracle Converged Database 23c MongoDB API Live Demo.pdf** file

# Summary



## Modern applications need to generate value from data in new ways

- They are built using new development methodologies and technologies
- But it complicates Database Architecture by using multiple single-purpose databases
- Development focused on Integration

## Oracle Database allows for easy start developing document-oriented applications

- Provides multiple APIs
- Fully supports document-oriented data model based on JSON specification
- Provides MongoDB API for MongoDB applications
- Provides multiple enterprise-scale solutions, which increase its availability, security and performance

# Some Developer Resources for free

---

- [Oracle Database Express Edition \(Version 21c\)](#)
- [Pre-Built Oracle DB Developer VM](#)
- [Docker Support](#)
- Live SQL: without installation learn and share <http://livesql.oracle.com>)
- [Autonomous Database Free Trial](#)
- Free SQL courses:  
[Databases for Developers: Foundations](#)

